

1

Boston Housing

```
In [2]: 1 # Import necessary libraries
2 import numpy as np
3 import pandas as pd
4 from sklearn.datasets import load_boston
5 from sklearn.model_selection import train_test_split
6 from sklearn.preprocessing import StandardScaler
7 import matplotlib.pyplot as plt
8 import seaborn as sns
9
10 # Load the Boston Housing Prices dataset
11 boston = load_boston()
12
13 # Convert the dataset to a DataFrame
14 data = pd.DataFrame(data=np.c_[boston['data'], boston['target']], column
15
16 # Data Exploration
17 # Display the first few rows of the dataset
18 print(data.head())
19
20 # Summary statistics
21 print(data.describe())
22
23
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

	PTRATIO	B	LSTAT	target
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2

	CRIM	ZN	INDUS	CHAS	NOX	
RM \						
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000
000						
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284
634						
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702
617						
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561
000						
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885
500						
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208
500						
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623
500						
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780
000						

	AGE	DIS	RAD	TAX	PTRATIO	
B \						
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000
000						
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674
032						
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294
864						
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320
000						
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377
500						
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440
000						
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225
000						
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900
000						

	LSTAT	target
count	506.000000	506.000000
mean	12.653063	22.532806
std	7.141062	9.197104
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.360000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000

C:\Users\Anusha V\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function load_boston is deprecated; `load_boston` is deprecated in 1.0 and will be removed in 1.2.

The Boston housing prices dataset has an ethical problem. You can refer to the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```
import pandas as pd
import numpy as np
```

```
data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
e) data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]
```

Alternative datasets include the California housing dataset (i.e. :func:`~sklearn.datasets.fetch_california_housing`) and the Ames housing dataset. You can load the datasets as follows::

```
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
```

for the California housing dataset and::

```
from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)
```

for the Ames housing dataset.

```
warnings.warn(msg, category=FutureWarning)
```

```
In [4]: 1 # Import necessary libraries
2 from sklearn.model_selection import train_test_split
3
4 # Split the data into training and testing sets
5 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
6
7 # Check the shape of the split datasets
8 print("X_train shape:", X_train.shape)
9 print("X_test shape:", X_test.shape)
10 print("y_train shape:", y_train.shape)
11 print("y_test shape:", y_test.shape)
12
```

X_train shape: (404, 13)

X_test shape: (102, 13)


y_train shape: (404,)

y_test shape: (102,)

```

In [3]: 1 # Correlation matrix
2 correlation_matrix = data.corr()
3 plt.figure(figsize=(12, 8))
4 sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
5 plt.title('Correlation Matrix')
6 plt.show()
7
8 # Data Preprocessing
9 # Split the data into features (X) and target (y)
10 X = data.drop('target', axis=1)
11 y = data['target']
12
13 # Standardize the features (mean=0, std=1)
14 scaler = StandardScaler()
15 X = scaler.fit_transform(X)
16
17 # Data Splitting
18 # Split the data into training and testing sets
19 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
20
21 # Check the shape of the split datasets
22 print("X_train shape:", X_train.shape)
23 print("X_test shape:", X_test.shape)
24 print("y_train shape:", y_train.shape)
25 print("y_test shape:", y_test.shape)
26

```



X_train shape: (404, 13)

X_test shape: (102, 13)

y_train shape: (404,)

y_test shape: (102,)

```
In [9]: 1  ##Evaluation matrix
2
3  ##Import necessary Libraries
4  from sklearn.linear_model import LinearRegression
5  from sklearn.metrics import mean_squared_error, r2_score
6
7  # Create a LinearRegression model
8  model = LinearRegression()
9
10 # Fit the model on the training data
11 model.fit(X_train[:, 5:6], y_train) # Using the 6th feature as an exam
12
13 # Make predictions on the test data
14 y_pred = model.predict(X_test[:, 5:6])
15
16 # Evaluate the model
17 mse = mean_squared_error(y_test, y_pred)
18 r2 = r2_score(y_test, y_pred)
19
20 # Print the model's coefficients and performance metrics
21 print("Coefficients:", model.coef_)
22 print("Intercept:", model.intercept_)
23 print("Mean Squared Error (MSE):", mse)
24 print("R-squared (R2) Score:", r2)
25
26 from sklearn.metrics import mean_absolute_error
27
28 # Calculate the Mean Absolute Error (MAE)
29 mae = mean_absolute_error(y_test, y_pred)
30
31 # Print the MAE
32 print("Mean Absolute Error (MAE):", mae)
33
```

Coefficients: [6.56178323]
Intercept: 22.50433758446674
Mean Squared Error (MSE): 46.144775347317264
R-squared (R2) Score: 0.3707569232254778
Mean Absolute Error (MAE): 4.478335832064148

```
In [16]: 1  from sklearn.model_selection import cross_val_score
2  from sklearn.linear_model import LinearRegression
3
4  # Create a LinearRegression model
5  model = LinearRegression()
6
7  # Perform cross-validation with 5 folds
8  scores = cross_val_score(model, X, y, cv=5, scoring='neg_mean_squared_e
9
10 # Calculate the mean of the negative mean squared error scores
11 mean_mse = -scores.mean()
12
13 # Print the mean squared error
14 print("Mean Squared Error (MSE) for Cross-Validation:", mean_mse)
15
```

Mean Squared Error (MSE) for Cross-Validation: 37.13180746769907

```
In [8]: 1  ##Holdout Method (Train-Test Split):
2
3  from sklearn.model_selection import train_test_split
4  from sklearn.metrics import mean_squared_error
5
6  # Split the data into training and testing sets
7  X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2
8
9  # Create and fit the model
10 model = LinearRegression()
11 model.fit(X_train[:, 5:6], y_train)
12
13 # Make predictions on the test data
14 y_pred = model.predict(X_test[:, 5:6])
15
16 # Evaluate the model using mean squared error
17 mse = mean_squared_error(y_test, y_pred)
18 print("Holdout Method - Mean Squared Error (MSE):", mse)
19
```

Holdout Method - Mean Squared Error (MSE): 46.144775347317264


```
In [11]: 1  ##Leave-One-Out Cross-Validation (LOOCV):
2
3  from sklearn.model_selection import LeaveOneOut
4
5  loo = LeaveOneOut()
6  mse_list = []
7
8  for train_index, test_index in loo.split(X):
9      X_train, X_test = X[train_index], X[test_index]
10     y_train, y_test = y[train_index], y[test_index]
11
12     model = LinearRegression()
13     model.fit(X_train[:, 5:6], y_train)
14
15     y_pred = model.predict(X_test[:, 5:6])
16     mse = mean_squared_error(y_test, y_pred)
17     mse_list.append(mse)
18
19 average_mse_loocv = np.mean(mse_list)
20 print("Leave-One-Out Cross-Validation (LOOCV) - Average MSE:", average_
21
```

Leave-One-Out Cross-Validation (LOOCV) - Average MSE: 44.216664193967986


```
In [12]: 1  ##K-Fold Cross-Validation (e.g., K=5):
2
3  from sklearn.model_selection import KFold
4
5  k = 5
6  kf = KFold(n_splits=k, shuffle=True, random_state=42)
7  mse_list = []
8
9  for train_index, test_index in kf.split(X):
10     X_train, X_test = X[train_index], X[test_index]
11     y_train, y_test = y[train_index], y[test_index]
12
13     model = LinearRegression()
14     model.fit(X_train[:, 5:6], y_train)
15
16     y_pred = model.predict(X_test[:, 5:6])
17     mse = mean_squared_error(y_test, y_pred)
18     mse_list.append(mse)
19
20 average_mse_kfold = np.mean(mse_list)
21 print(f"{k}-Fold Cross-Validation - Average MSE:", average_mse_kfold)
22
```

5-Fold Cross-Validation - Average MSE: 43.84204126558527

```
In [14]: 1  ##Simple Liner Regression
2
3
4  import matplotlib.pyplot as plt
5  from sklearn.linear_model import LinearRegression
6
7  # Create a LinearRegression model
8  model = LinearRegression()
9
10 # Fit the model on the training data
11 model.fit(X_train[:, 5:6], y_train) # Using the 6th feature as an exam
12
13 # Make predictions on the entire dataset
14 y_pred = model.predict(X[:, 5:6])
15
16 # Create a scatter plot of the actual data points
17 plt.scatter(X[:, 5], y, color='blue', label='Actual Data')
18
19 # Overlay the regression line
20 plt.plot(X[:, 5], y_pred, color='red', linewidth=2, label='Regression L
21
22 # Set labels and title
23 plt.xlabel('Feature')
24 plt.ylabel('Target')
25 plt.title('Simple Linear Regression')
26
27 # Add a Legend
28 plt.legend()
29
30 # Show the plot
31 plt.show()
32
```



In [16]:

```
1  ##multiple linear regression
2
3  import numpy as np
4  import pandas as pd
5  from sklearn.datasets import load_boston
6  from sklearn.model_selection import train_test_split
7  from sklearn.linear_model import LinearRegression
8  import matplotlib.pyplot as plt
9
10 # Load the Boston Housing Prices dataset
11 boston = load_boston()
12 data = pd.DataFrame(data=np.c_[boston['data'], boston['target']], columns=['data', 'target'])
13
14 # Split the data into features (X) and target (y)
15 X = data.drop('target', axis=1)
16 y = data['target']
17
18 # Split the data into training and testing sets
19 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
20
21 # Create a Linear Regression model
22 model = LinearRegression()
23
24 # Fit the model on the training data
25 model.fit(X_train, y_train)
26
27 # Make predictions on the test data
28 y_pred = model.predict(X_test)
29
30 # Visualize the results
31 plt.scatter(y_test, y_pred)
32 plt.xlabel("Actual Prices")
33 plt.ylabel("Predicted Prices")
34 plt.title("Multiple Linear Regression")
35 plt.show()
36
```

```
C:\Users\Anusha V\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function load_boston is deprecated; `load_boston` is deprecated in 1.0 and will be removed in 1.2.
```

The Boston housing prices dataset has an ethical problem. You can refer to the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```
import pandas as pd
import numpy as np

data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
e) data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]
```

Alternative datasets include the California housing dataset (i.e. :func:`~sklearn.datasets.fetch_california_housing`) and the Ames housing dataset. You can load the datasets as follows::


```
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()

for the California housing dataset and::

from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)

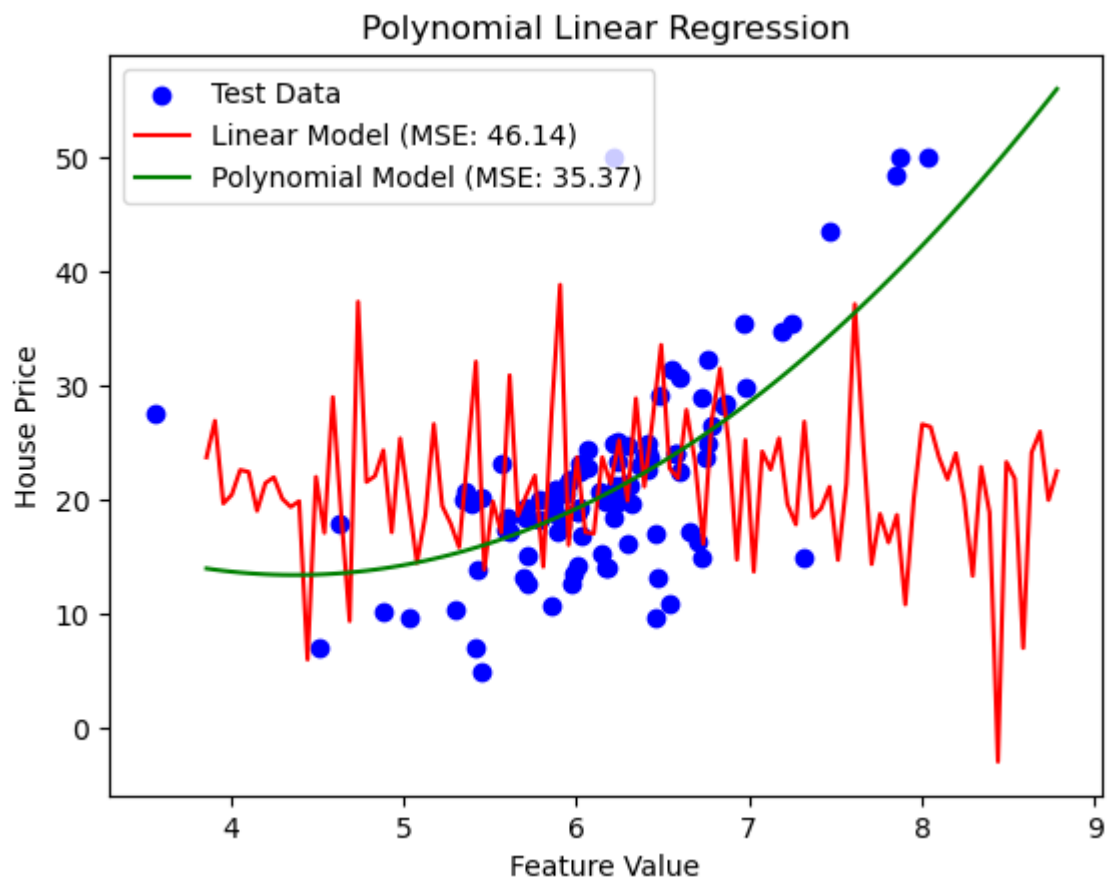
for the Ames housing dataset.
```

```
warnings.warn(msg, category=FutureWarning)
```



In [24]:

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.linear_model import LinearRegression
4 from sklearn.preprocessing import PolynomialFeatures
5 from sklearn.metrics import mean_squared_error
6 from sklearn.model_selection import train_test_split
7
8 # Split the data into training and testing sets
9 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
10
11 # Create a LinearRegression model
12 linear_model = LinearRegression()
13
14 # Select a feature for polynomial regression (e.g., the 6th feature)
15 feature_idx = 5
16
17 # Extract the chosen feature as a Pandas Series
18 X_train_feature = X_train.iloc[:, feature_idx].values.reshape(-1, 1)
19 X_test_feature = X_test.iloc[:, feature_idx].values.reshape(-1, 1)
20
21 # Ensure that X_train_feature and y_train have the same number of data
22 X_train_feature = X_train_feature[:len(y_train)]
23
24 # Fit the linear model
25 linear_model.fit(X_train_feature, y_train)
26
27 # Generate predictions for the linear model
28 y_pred_linear = linear_model.predict(X_test_feature)
29
30 # Create a range of values for the x-axis (for plotting)
31 x_range = np.linspace(X_train_feature.min(), X_train_feature.max(), num
32
33 # Create a PolynomialFeatures transformer (e.g., quadratic degree=2)
34 poly = PolynomialFeatures(degree=2)
35 X_poly = poly.fit_transform(X_train_feature)
36
37 # Fit a polynomial regression model
38 poly_model = LinearRegression()
39 poly_model.fit(X_poly, y_train)
40
41 # Generate predictions for the polynomial model
42 X_poly_range = poly.transform(x_range)
43 y_pred_poly = poly_model.predict(X_poly_range)
44
45 # Calculate the Mean Squared Error (MSE) for both models
46 mse_linear = mean_squared_error(y_test, y_pred_linear)
47 mse_poly = mean_squared_error(y_test, poly_model.predict(poly.transform
48
49 # Plot the results
50 plt.scatter(X_test_feature, y_test, label='Test Data', color='b')
51 plt.plot(x_range, y_pred_linear, label=f'Linear Model (MSE: {mse_linear
52 plt.plot(x_range, y_pred_poly, label=f'Polynomial Model (MSE: {mse_poly
53 plt.xlabel("Feature Value")
54 plt.ylabel("House Price")
55 plt.legend()
56 plt.title("Polynomial Linear Regression")
57 plt.show()
```



Build a SVM Model in python for Fish

dataset from Kaggl


```
In [2]: 1 import pandas as pd
2 from sklearn.model_selection import train_test_split
3 from sklearn.svm import SVC
4 from sklearn.preprocessing import StandardScaler
5 from sklearn.metrics import accuracy_score
6
7 data = pd.read_csv(r"C:\Users\Anusha V\Downloads\fish.csv")
8 X = data[['nofish', 'livebait', 'camper']]
9 y = data['persons']
10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
11 scaler = StandardScaler()
12 X_train = scaler.fit_transform(X_train)
13 X_test = scaler.transform(X_test)
14 svm_model = SVC()
15 svm_model.fit(X_train, y_train)
16 y_pred = svm_model.predict(X_test)
17 accuracy = accuracy_score(y_test, y_pred)
18 print(f"Accuracy of the SVM model: {accuracy}")
19
```

```
C:\Users\Anusha V\anaconda3\lib\site-packages\scipy\__init__.py:155: UserWarning: A NumPy version >=1.18.5 and <1.25.0 is required for this version of SciPy (detected version 1.26.1)
  warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

Accuracy of the SVM model: 0.28

gini and entropy

The Gini impurity is a measure of how often a randomly chosen element from a set would be incorrectly labeled if it was randomly labeled according to the distribution of labels in the set.

```
In [3]: 1 import pandas as pd
2
3 y = data['camper']
4 def gini_impurity(labels):
5     total_samples = len(labels)
6     label_counts = labels.value_counts()
7     impurity = 1
8     for count in label_counts:
9         label_probability = count / total_samples
10        impurity -= label_probability ** 2
11    return impurity
12 gini = gini_impurity(y)
13 print(f"Gini Impurity for '{X}': {gini}")
14
```

```
Gini Impurity for '    nofish  livebait  camper
0         1         0         0
1         0         1         1
2         0         1         0
3         0         1         1
4         0         1         0
..      ...      ...      ...
245        1         1         1
246        0         1         1
247        0         1         1
248        1         1         1
249        1         1         1
```

```
[250 rows x 3 columns]': 0.48451200000000005
```

```
In [4]: 1 import pandas as pd
2 import math
3
4 y = data['camper']
5 def entropy(labels):
6     total_samples = len(y)
7     label_counts = y.value_counts()
8     entropy_val = 0
9     for count in label_counts:
10        label_probability = count / total_samples
11        entropy_val -= label_probability * math.log(label_probability,
12    return entropy_val
13 entropy_result = entropy(y)
14 print(f"Entropy for '{y}': {entropy_result}")
15
```

```
Entropy for '0      0
1      1
2      0
3      1
4      0
..
245    1
246    1
247    1
248    1
249    1
```

```
Name: camper, Length: 250, dtype: int64': 0.9775387286988189
```


Random Forest Classifier:

```
In [10]: 1 from sklearn.datasets import load_iris
2 from sklearn.ensemble import RandomForestClassifier
3 from sklearn.model_selection import train_test_split
4 from sklearn.metrics import accuracy_score
5 data = load_iris()
6 X = data.data
7 y = data.target
8 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
9 clf = RandomForestClassifier(n_estimators=100, random_state=42)
10 clf.fit(X_train, y_train)
11 predictions = clf.predict(X_test)
12 accuracy = accuracy_score(y_test, predictions)
13 print(f"Accuracy of Random Forest Classifier: {accuracy}")
```

Accuracy of Random Forest Classifier: 1.0


K-means Clustering

```
In [9]: 1 import pandas as pd
2 from sklearn.cluster import KMeans
3 from sklearn.preprocessing import StandardScaler
4 import matplotlib.pyplot as plt
5
6
7 features = data[['nofish', 'camper']]
8 scaler = StandardScaler()
9 scaled_features = scaler.fit_transform(features)
10 num_clusters = 3 # You can change this value as needed
11 kmeans = KMeans(n_clusters=num_clusters)
12 kmeans.fit(scaled_features)
13 data['cluster'] = kmeans.labels_
14 plt.scatter(data['nofish'], data['camper'], c=data['cluster'], cmap='viridis')
15 plt.title('K-means Clustering')
16 plt.xlabel('Feature 1')
17 plt.ylabel('Feature 2')
18 plt.show()
19
```



PCA

```
In [19]: 1 import pandas as pd
2 from sklearn.decomposition import PCA
3 from sklearn.preprocessing import StandardScaler
4 import matplotlib.pyplot as plt
5
6 features = data.drop(columns=['nofish', 'livebait'])
7 # Standardize the features
8 scaler = StandardScaler()
9 scaled_features = scaler.fit_transform(features)
10 # Apply PCA
11 pca = PCA(n_components=2) # Choose the number of components you want to
12 principal_components = pca.fit_transform(scaled_features)
13 # Create a DataFrame for the principal components
14 principal_df = pd.DataFrame(data=principal_components, columns=['PC1',
15 # Visualize the PCA results
16 plt.figure(figsize=(8, 6))
17 plt.scatter(principal_df['PC1'], principal_df['PC2'], alpha=0.5)
18 plt.title('PCA: Two Principal Components')
19 plt.xlabel('Principal Component 1')
20 plt.ylabel('Principal Component 2')
21 plt.show()
22
```



MLOps


```
In [26]: 1 from sklearn.datasets import load_iris
2 from sklearn.model_selection import train_test_split
3 from sklearn.preprocessing import StandardScaler
4 from sklearn.decomposition import PCA
5 from sklearn.svm import SVC
6 from sklearn.pipeline import Pipeline
7 from sklearn.metrics import accuracy_score
8 data = load_iris()
9 X = data.data
10 y = data.target
11 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
12 steps = [
13     ('scaler', StandardScaler()), # Step 1: Standard Scaling
14     ('pca', PCA(n_components=2)), # Step 2: PCA for dimensionality reduction
15     ('classifier', SVC(kernel='rbf', C=1.0)) # Step 3: Support Vector Classification
16 ]
17 pipeline = Pipeline(steps)
18 pipeline.fit(X_train, y_train)
19 predictions = pipeline.predict(X_test)
20 accuracy = accuracy_score(y_test, predictions)
21 print(f"Accuracy of the Pipeline: {accuracy}")
22
```

Accuracy of the Pipeline: 0.9333333333333333

Basic Ensemble Techniques

Max Voting:

Max Voting is one of the simplest ensemble techniques where you combine the predictions of multiple models and choose the majority class as the final prediction.

```
In [35]: 1 from sklearn.datasets import load_iris # Import the dataset you need
2
3 # Load the dataset (for example, the Iris dataset)
4 data = load_iris()
5 X = data.data # Features
6 y = data.target # Labels
7
```

```
In [36]: 1 from sklearn.ensemble import VotingClassifier
2 from sklearn.model_selection import train_test_split
3 from sklearn.linear_model import LogisticRegression
4 from sklearn.tree import DecisionTreeClassifier
5 from sklearn.svm import SVC
6
7 # Create individual models
8 model1 = LogisticRegression()
9 model2 = DecisionTreeClassifier()
10 model3 = SVC()
11
12 # Create a Voting Classifier
13 ensemble_model = VotingClassifier(estimators=[('lr', model1), ('dt', mo
14
15 # Load your dataset and split it into train and test sets
16
17 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3
18
19 # Fit the ensemble model on the training data
20 ensemble_model.fit(X_train, y_train)
21
22 # Make predictions using the ensemble model
23 y_pred = ensemble_model.predict(X_test)
24 y_pred
```

```
Out[36]: array([1, 0, 0, 2, 2, 2, 1, 2, 0, 0, 1, 0, 2, 2, 1, 1, 0, 1, 0, 1, 1, 2,
                2, 1, 2, 0, 2, 1, 0, 1, 0, 0, 0, 2, 2, 1, 0, 1, 0, 2, 1, 2, 1, 0,
                1])
```

Averaging:

Averaging is an ensemble technique where you average the predictions of multiple models to get the final prediction.


```

In [37]: 1 from sklearn.model_selection import train_test_split
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.tree import DecisionTreeClassifier
4 from sklearn.svm import SVC
5
6 # Create individual models
7 model1 = LogisticRegression()
8 model2 = DecisionTreeClassifier()
9 model3 = SVC()
10
11 # Load your dataset and split it into train and test sets
12 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
13
14 # Fit each model on the training data
15 model1.fit(X_train, y_train)
16 model2.fit(X_train, y_train)
17 model3.fit(X_train, y_train)
18
19 # Make predictions using each model
20 pred1 = model1.predict(X_test)
21 pred2 = model2.predict(X_test)
22 pred3 = model3.predict(X_test)
23
24 # Average the predictions
25 final_pred = (pred1 + pred2 + pred3) / 3
26 final_pred

```

```

Out[37]: array([[1.          , 1.          , 0.          , 0.          , 0.          ,
1.          , 2.          , 2.          , 0.          , 0.          ,
0.          , 0.          , 1.          , 2.          , 1.          ,
0.          , 1.          , 1.33333333, 1.          , 1.          ,
1.          , 0.          , 0.          , 1.33333333, 1.66666667,
0.          , 1.          , 1.          , 0.          , 0.          ,
0.          , 1.          , 1.          , 2.          , 2.          ,
2.          , 0.          , 0.          , 1.          , 0.          ,
0.          , 1.66666667, 2.          , 2.          , 1.          ]])

```

Weighted Average:

Weighted Averaging is similar to Averaging, but you assign different weights to the models

```
In [38]: 1 from sklearn.model_selection import train_test_split
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.tree import DecisionTreeClassifier
4 from sklearn.svm import SVC
5
6 # Create individual models
7 model1 = LogisticRegression()
8 model2 = DecisionTreeClassifier()
9 model3 = SVC()
10
11 # Load your dataset and split it into train and test sets
12 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
13
14 # Fit each model on the training data
15 model1.fit(X_train, y_train)
16 model2.fit(X_train, y_train)
17 model3.fit(X_train, y_train)
18
19 # Make predictions using each model
20 pred1 = model1.predict(X_test)
21 pred2 = model2.predict(X_test)
22 pred3 = model3.predict(X_test)
23
24 # Assign weights to each model
25 weight1 = 0.3
26 weight2 = 0.4
27 weight3 = 0.3
28
29 # Calculate the weighted average
30 final_pred = (weight1 * pred1 + weight2 * pred2 + weight3 * pred3)
31 final_pred
```

```
Out[38]: array([2., 2., 0., 0., 1., 2., 2., 1., 2., 1., 2., 2., 2., 0., 0., 0., 2.,
                2., 2., 2., 2., 2., 2., 1., 0., 0., 1., 1., 1., 2., 0., 1., 0., 2.,
                0., 1., 1., 0., 1., 1., 2., 0., 1., 1., 0.])
```

Advanced Ensemble Techniques

Stacking:

Stacking is an ensemble technique that combines the predictions of multiple models using another model, called a meta-learner. Here's a simple example of stacking using scikit-learn:

```
In [39]: 1 from sklearn.ensemble import StackingClassifier
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.ensemble import RandomForestClassifier
4 from sklearn.svm import SVC
5 from sklearn.datasets import load_iris
6 from sklearn.model_selection import train_test_split
7
8 # Load a sample dataset
9 data = load_iris()
10 X, y = data.data, data.target
11
12 # Split the data into a training set and a test set
13 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
14
15 # Base models
16 base_models = [
17     ('rf', RandomForestClassifier(n_estimators=100, random_state=42)),
18     ('svc', SVC(kernel='linear', C=1, probability=True))
19 ]
20
21 # Meta-Learner
22 meta_learner = LogisticRegression()
23
24 # Create the stacking classifier
25 stacking_model = StackingClassifier(estimators=base_models, final_estimator=meta_learner)
26
27 # Fit the stacking model on the training data
28 stacking_model.fit(X_train, y_train)
29
30 # Make predictions using the stacked model
31 y_pred = stacking_model.predict(X_test)
32
33 # Evaluate the performance of the stacking model
34 from sklearn.metrics import accuracy_score
35 accuracy = accuracy_score(y_test, y_pred)
36 print("Accuracy:", accuracy)
37
```

Accuracy: 1.0

Blending:

Blending is similar to stacking but uses a separate validation set to create meta-features for the final prediction. Here's a simple example of blending:

```
In [40]: 1 from sklearn.svm import SVC
2 from sklearn.datasets import load_iris
3 from sklearn.model_selection import train_test_split
4
5 # Load a sample dataset
6 data = load_iris()
7 X, y = data.data, data.target
8
9 # Split the data into a training set and a test set
10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
11
12 # Train base models
13 rf = RandomForestClassifier(n_estimators=100, random_state=42)
14 svc = SVC(kernel='linear', C=1, probability=True)
15
16 rf.fit(X_train, y_train)
17 svc.fit(X_train, y_train)
18
19 # Create meta-features using validation set
20 rf_preds = rf.predict(X_test)
21 svc_preds = svc.predict(X_test)
22
23 # Combine predictions
24 blend_preds = (rf_preds + svc_preds) / 2
25
26 # Evaluate the performance of the blending model
27 from sklearn.metrics import accuracy_score
28 accuracy = accuracy_score(y_test, blend_preds)
29 print("Accuracy:", accuracy)
30
```

Accuracy: 1.0

Bagging:

Bagging is an ensemble technique that uses bootstrap resampling to train multiple base models. Here's a simple example using the RandomForestClassifier in scikit-learn:

```
In [41]: 1 from sklearn.ensemble import RandomForestClassifier
2 from sklearn.datasets import load_iris
3 from sklearn.model_selection import train_test_split
4
5 # Load a sample dataset
6 data = load_iris()
7 X, y = data.data, data.target
8
9 # Split the data into a training set and a test set
10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
11
12 # Create a bagging classifier
13 bagging_model = RandomForestClassifier(n_estimators=100, random_state=42)
14
15 # Fit the bagging model on the training data
16 bagging_model.fit(X_train, y_train)
17
18 # Make predictions using the bagging model
19 y_pred = bagging_model.predict(X_test)
20
21 # Evaluate the performance of the bagging model
22 from sklearn.metrics import accuracy_score
23 accuracy = accuracy_score(y_test, y_pred)
24 print("Accuracy:", accuracy)
25
```

Accuracy: 1.0

Random forest regression¶

```
In [42]: 1 from sklearn.ensemble import RandomForestRegressor
2 from sklearn.datasets import make_regression
3 X, y = make_regression(n_features=4, n_informative=2, random_state=0, s
4 rfr = RandomForestRegressor(max_depth=3)
5 rfr.fit(X, y)
6 print(rfr.predict([[0, 1, 0, 1]]))
7
```

[32.55050435]

LogisticRegression

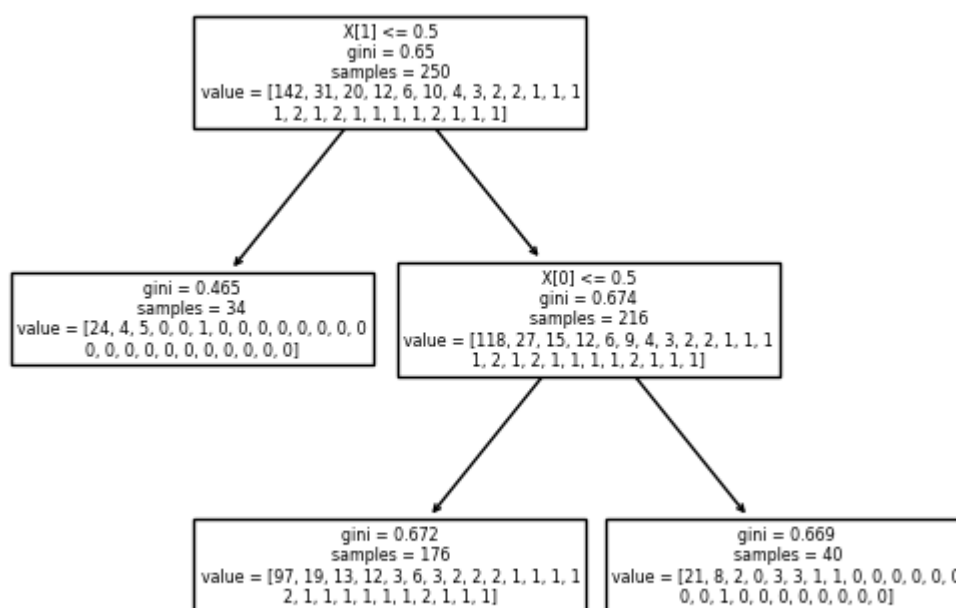
```
In [59]: 1 import pandas as pd
2 from sklearn.model_selection import train_test_split
3 from sklearn.linear_model import LogisticRegression
4 from sklearn.metrics import accuracy_score
5 import matplotlib.pyplot as plt
6
7 X = data.drop('nofish', axis=1)
8 y = data['livebait']
9 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
10 model = LogisticRegression()
11 model.fit(X_train, y_train)
12 y_pred = model.predict(X_test)
13 accuracy = accuracy_score(y_test, y_pred)
14 print(f"Accuracy: {accuracy}")
15
```

Accuracy: 1.0

DecisionTreeClassifier

```
In [60]: 1 import pandas as pd
2
3 X = data.iloc[:, :-1]
4 from sklearn.tree import DecisionTreeClassifier
5 clf1 = DecisionTreeClassifier()
6 clf1.fit(data.iloc[:, 0:2], data.iloc[:, -1])
7 from sklearn.tree import plot_tree
8 plot_tree(clf1)
9
```

```
Out[60]: [Text(0.4, 0.8333333333333334, 'X[1] <= 0.5\ngini = 0.65\nsamples = 250\nvalue = [142, 31, 20, 12, 6, 10, 4, 3, 2, 2, 1, 1, 1\n1, 2, 1, 2, 1, 1, 1, 1]'),
Text(0.2, 0.5, 'gini = 0.465\nsamples = 34\nvalue = [24, 4, 5, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]'),
Text(0.6, 0.5, 'X[0] <= 0.5\ngini = 0.674\nsamples = 216\nvalue = [118, 27, 15, 12, 6, 9, 4, 3, 2, 2, 1, 1, 1\n1, 2, 1, 2, 1, 1, 1, 1, 2, 1, 1, 1]'),
Text(0.4, 0.16666666666666666, 'gini = 0.672\nsamples = 176\nvalue = [97, 19, 13, 12, 3, 6, 3, 2, 2, 2, 1, 1, 1, 1\n2, 1, 1, 1, 1, 1, 1, 2, 1, 1, 1]'),
Text(0.8, 0.16666666666666666, 'gini = 0.669\nsamples = 40\nvalue = [21, 8, 2, 0, 3, 3, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]')]
```




CRICKET MATCH RESULT-PAST DATA

```
In [1]: 1 import pandas as pd
        2 data = pd.read_csv(r"C:\Users\Anusha V\Desktop\cricket.csv")
        3 data
```

```
Out[1]:
```

	id	season	city	date	team1	team2	toss_winner	toss_decision	re
0	1.0	2017.0	Hyderabad	05-04-2017	Sunrisers Hyderabad	Royal Challengers Bangalore	Royal Challengers Bangalore	field	nor
1	2.0	2017.0	Pune	06-04-2017	Mumbai Indians	Rising Pune Supergiant	Rising Pune Supergiant	field	nor
2	3.0	2017.0	Rajkot	07-04-2017	Gujarat Lions	Kolkata Knight Riders	Kolkata Knight Riders	field	nor
3	4.0	2017.0	Indore	08-04-2017	Rising Pune Supergiant	Kings XI Punjab	Kings XI Punjab	field	nor
4	5.0	2017.0	Bangalore	08-04-2017	Royal Challengers Bangalore	Delhi Daredevils	Royal Challengers Bangalore	bat	nor
...
631	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	↑
632	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	↑
633	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	↑
634	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	↑
635	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	↑

636 rows × 18 columns



```
In [20]: 1 run_a_wins = len(data[data['win_by_runs'] == 'run A'])
        2 hs_b_wins = len(data[data['win_by_wickets'] == 'highest_scores B'])
        3 print("Number of matches won by Team A:", run_a_wins)
        4 print("Number of matches won by Team B:", hs_b_wins)
        5
```

Number of matches won by Team A: 0
Number of matches won by Team B: 0

```
In [45]: 1 import pandas as pd
        2 import numpy as np
        3 from sklearn.model_selection import train_test_split
        4 from sklearn.preprocessing import StandardScaler
        5
```



```
In [46]: 1  ##Data Exploration:
2
3  # Display the first few rows of the dataset
4  print(data.head())
5
6  # Summary statistics
7  print(data.describe())
8
9  # Check for missing values
10 print(data.isnull().sum())
11
12 # Visualize data as needed (e.g., using matplotlib or seaborn)
13
```

Empty DataFrame

Columns: [id, season, city, date, team1, team2, toss_winner, toss_decision, result, dl_applied, winner, win_by_runs, win_by_wickets, player_of_match, venue, umpire1, umpire2, umpire3]

Index: []

	id	season	dl_applied	win_by_runs	win_by_wickets	umpire3
count	0.0	0.0	0.0	0.0	0.0	0.0
mean	NaN	NaN	NaN	NaN	NaN	NaN
std	NaN	NaN	NaN	NaN	NaN	NaN
min	NaN	NaN	NaN	NaN	NaN	NaN
25%	NaN	NaN	NaN	NaN	NaN	NaN
50%	NaN	NaN	NaN	NaN	NaN	NaN
75%	NaN	NaN	NaN	NaN	NaN	NaN
max	NaN	NaN	NaN	NaN	NaN	NaN
id			0.0			
season			0.0			
city			0.0			
date			0.0			
team1			0.0			
team2			0.0			
toss_winner			0.0			
toss_decision			0.0			
result			0.0			
dl_applied			0.0			
winner			0.0			
win_by_runs			0.0			
win_by_wickets			0.0			
player_of_match			0.0			
venue			0.0			
umpire1			0.0			
umpire2			0.0			
umpire3			0.0			

dtype: float64

```
In [53]: 1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.linear_model import LinearRegression
4 from sklearn.preprocessing import PolynomialFeatures
5 from sklearn.metrics import mean_squared_error
6 from sklearn.model_selection import train_test_split
7 from sklearn.preprocessing import StandardScaler
8
9 # Data Exploration
10 # (Insert code to explore the dataset)
11
12 # Handle missing values (if any)
13 data.dropna(inplace=True)
14
15 # Encode categorical variables (if needed) using techniques like one-hot
16 # (Insert code to encode categorical variables)
17
18 # Data Preprocessing
19
20 # Extract the target variable (predicting the winning team)
21 y = data['winner']
22
23 # Extract relevant features (predictor variables)
24 # Here, you can select the features that you believe are relevant
25 X = data[['team1', 'team2', 'toss_winner', 'toss_decision', 'venue']]
26
27 # Check if there are enough samples for splitting
28 if len(data) > 0:
29     # Data Splitting
30     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=
31
32     # Check the shape of the split datasets
33     print("X_train shape:", X_train.shape)
34     print("X_test shape:", X_test.shape)
35     print("y_train shape:", y_train.shape)
36     print("y_test shape:", y_test.shape)
37
38     # Standardize/normalize the features
39     scaler = StandardScaler()
40     X_train = scaler.fit_transform(X_train)
41     X_test = scaler.transform(X_test)
42
43     # Continue with the rest of your analysis or modeling.
44
```

In [55]:

1X_train

Out[55]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	
477	15.02340	0.0	18.10	0.0	0.6140	5.304	97.3	2.1007	24.0	666.0	20.2	349.4
15	0.62739	0.0	8.14	0.0	0.5380	5.834	56.5	4.4986	4.0	307.0	21.0	395.6
332	0.03466	35.0	6.06	0.0	0.4379	6.031	23.3	6.6407	1.0	304.0	16.9	362.2
423	7.05042	0.0	18.10	0.0	0.6140	6.103	85.1	2.0218	24.0	666.0	20.2	2.5
19	0.72580	0.0	8.14	0.0	0.5380	5.727	69.5	3.7965	4.0	307.0	21.0	390.9
...
106	0.17120	0.0	8.56	0.0	0.5200	5.836	91.9	2.2110	5.0	384.0	20.9	395.6
270	0.29916	20.0	6.96	0.0	0.4640	5.856	42.1	4.4290	3.0	223.0	18.6	388.6
348	0.01501	80.0	2.01	0.0	0.4350	6.635	29.7	8.3440	4.0	280.0	17.0	390.9
435	11.16040	0.0	18.10	0.0	0.7400	6.629	94.6	2.1247	24.0	666.0	20.2	109.8
102	0.22876	0.0	8.56	0.0	0.5200	6.405	85.4	2.7147	5.0	384.0	20.9	70.8

404 rows × 13 columns

In [56]:

1X_test


Out[56]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	
173	0.09178	0.0	4.05	0.0	0.510	6.416	84.1	2.6463	5.0	296.0	16.6	395.5
274	0.05644	40.0	6.41	1.0	0.447	6.758	32.9	4.0776	4.0	254.0	17.6	396.9
491	0.10574	0.0	27.74	0.0	0.609	5.983	98.8	1.8681	4.0	711.0	20.1	390.1
72	0.09164	0.0	10.81	0.0	0.413	6.065	7.8	5.2873	4.0	305.0	19.2	390.9
452	5.09017	0.0	18.10	0.0	0.713	6.297	91.8	2.3682	24.0	666.0	20.2	385.0
...
412	18.81100	0.0	18.10	0.0	0.597	4.628	100.0	1.5539	24.0	666.0	20.2	28.7
436	14.42080	0.0	18.10	0.0	0.740	6.461	93.3	2.0026	24.0	666.0	20.2	27.4
411	14.05070	0.0	18.10	0.0	0.597	6.657	100.0	1.5275	24.0	666.0	20.2	35.0
86	0.05188	0.0	4.49	0.0	0.449	6.015	45.1	4.4272	3.0	247.0	18.5	395.9
75	0.09512	0.0	12.83	0.0	0.437	6.286	45.0	4.5026	5.0	398.0	18.7	383.2

102 rows × 13 columns

```
In [4]: 1 import matplotlib.pyplot as plt
        2 plt.scatter(data['win_by_wickets'], data['win_by_runs'])
```

Out[4]: <matplotlib.collections.PathCollection at 0x13fd2f1b490>



```
In [ ]: 1
```

```
In [ ]: 1
```

CROP YIELD - PAST DATA

```
In [10]: 1 import pandas as pd
2
3 # Replace the file path with the correct one
4 file_path = r"C:\Users\Anusha V\Desktop\CropYield.xlsx"
5
6 try:
7     data1 = pd.read_excel(file_path)
8 except Exception as e:
9     print(f"An error occurred: {str(e)}")
10
11 # Now you can work with the 'data1' DataFrame
12 print(data1)
13
```

		State	District	Crop	Crop_Year	\
0	Andaman and Nicobar Island	NICOBARS	Arecanut	2007		
1	Andaman and Nicobar Island	NICOBARS	Arecanut	2007		
2	Andaman and Nicobar Island	NICOBARS	Arecanut	2008		
3	Andaman and Nicobar Island	NICOBARS	Arecanut	2008		
4	Andaman and Nicobar Island	NICOBARS	Arecanut	2009		
..		
197	Andaman and Nicobar Island	SOUTH ANDAMANS	Cashewnut	2010		
198	Andaman and Nicobar Island	SOUTH ANDAMANS	Cashewnut	2011		
199	Andaman and Nicobar Island	SOUTH ANDAMANS	Cashewnut	2012		
200	Andaman and Nicobar Island	SOUTH ANDAMANS	Cashewnut	2013		
201	Andaman and Nicobar Island	SOUTH ANDAMANS	Cashewnut	2014		

	Season	Area	Production	Yield
0	Kharif	2439.6	3415.0	1.40
1	Rabi	1626.4	2277.0	1.40
2	Autumn	4147.0	3060.0	0.74
3	Summer	4147.0	2660.0	0.64
4	Autumn	4153.0	3120.0	0.75
..
197	Rabi	15.0	11.0	0.70
198	Rabi	20.0	15.0	0.77
199	Rabi	35.4	21.0	0.60
200	Rabi	29.5	22.0	0.74
201	Rabi	14.9	23.0	1.53

[202 rows x 8 columns]

```
In [11]: 1 import pandas as pd
2
3 # Display the first few rows of the dataset
4 print(data1.head())
5
6 # Check for missing values
7 print(data1.isnull().sum())
8
9 # Summary statistics
10 print(data1.describe())
11
```

	State	District	Crop	Crop_Year	Season
0	Andaman and Nicobar	Island	NICOBARS	Arecanut	2007 Kharif
1	Andaman and Nicobar	Island	NICOBARS	Arecanut	2007 Rabi
2	Andaman and Nicobar	Island	NICOBARS	Arecanut	2008 Autumn
3	Andaman and Nicobar	Island	NICOBARS	Arecanut	2008 Summer
4	Andaman and Nicobar	Island	NICOBARS	Arecanut	2009 Autumn

	Area	Production	Yield
0	2439.6	3415.0	1.40
1	1626.4	2277.0	1.40
2	4147.0	3060.0	0.74
3	4147.0	2660.0	0.64
4	4153.0	3120.0	0.75


```
State      0
District   0
Crop        0
Crop_Year  0
Season      0
Area        0
Production  1
Yield       0
dtype: int64
```


	Crop_Year	Area	Production	Yield
count	202.000000	202.000000	201.000000	202.000000
mean	2010.747525	746.185485	2220.756219	2.943020
std	5.393991	880.050543	3273.402037	4.637631
min	2000.000000	0.500000	0.000000	0.000000
25%	2006.000000	81.000000	21.000000	0.220000
50%	2011.000000	450.000000	208.000000	0.700000
75%	2015.000000	1034.300000	3602.000000	3.350000
max	2019.000000	4153.000000	17374.000000	26.450000

```
In [12]: 1 from sklearn.model_selection import train_test_split
2 import pandas as pd
3 data1 = pd.read_csv(r"C:\Users\Anusha V\Downloads\archive (8)\APY.csv")
4 X = data1.drop(columns=["Crop"])
5 y = data1["Crop"]
6 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
7
```

In [13]:

1 X_train

Out[13]:

	State	District	Crop_Year	Season	Area	Production	Yield
114808	Jharkhand	RAMGARH	2013	Winter	1361.0	7544.0	5.54
109257	Jammu and Kashmir	REASI	2015	Kharif	13.0	9.0	0.70
101966	Haryana	GURGAON	2009	Whole Year	39.0	800.0	20.51
313553	Uttar Pradesh	JALAUN	2010	Kharif	9.0	3.0	0.33
270181	Tamil Nadu	MADURAI	2006	Whole Year	6006.0	656204.0	109.26
...
119879	Karnataka	GADAG	2018	Kharif	13.0	7.0	0.54
259178	Tamil Nadu	TIRUCHIRAPPALLI	2002	Whole Year	5991.0	84200000.0	14054.41
131932	Karnataka	CHIKKABALLAPURA	2017	Kharif	63.0	14.0	0.22
146867	Kerala	IDUKKI	2014	Whole Year	4.0	NaN	0.00
121958	Karnataka	BELAGAVI	2017	Rabi	279.0	113.0	0.41

276268 rows × 7 columns

In [14]:

1 X_test

Out[14]:

	State	District	Crop_Year	Season	Area	Production	Yield
212198	Nagaland	PHEK	2017	Kharif	180.0	3650.0	20.28
2202	Andhra Pradesh	RANGAREDDI	2005	Whole Year	40.0	836.0	20.90
36666	Assam	JORHAT	2011	Whole Year	74.0	244.0	3.30
110896	Jammu and Kashmir	UDHAMPUR	2006	Rabi	2833.0	2079.0	0.73
212200	Nagaland	PHEK	2019	Kharif	200.0	4052.0	20.26
...
162204	Madhya Pradesh	MANDSAUR	2010	Kharif	37232.0	50743.0	1.36
35448	Assam	GOLAGHAT	2002	Kharif	40.0	18.0	0.45
185511	Maharashtra	BULDHANA	1997	Kharif	20000.0	16400.0	0.82
43783	Bihar	NALANDA	2017	Rabi	5369.0	6513.0	1.21
285273	Uttar Pradesh	ETAWAH	2019	Rabi	1566.0	5382.0	3.44

69068 rows × 7 columns

```
In [15]: 1 y_train  
        2
```

```
Out[15]: 114808          Potato  
        109257    Moong(Green Gram)  
        101966      Sweet potato  
        313553          Sannhamp  
        270181      Sugarcane  
        ...  
        119879      Castor seed  
        259178      Coconut  
        131932      Niger seed  
        146867          Onion  
        121958      Cowpea(Lobia)  
        Name: Crop, Length: 276268, dtype: object
```


```
In [16]: 1 y_test
```

```
Out[16]: 212198          Tapioca  
        2202          Banana  
        36666      Sweet potato  
        110896  Rapeseed &Mustard  
        212200          Tapioca  
        ...  
        162204          Maize  
        35448      Small millets  
        185511          Maize  
        43783          Gram  
        285273          Barley  
        Name: Crop, Length: 69068, dtype: object
```



```
In [17]: 1 import matplotlib.pyplot as plt
2 plt.scatter(data1['Crop_Year'], data1['Yield'])
```

Out[17]: <matplotlib.collections.PathCollection at 0x2ba2e4a0f70>



k mean

```
In [43]: 1 import numpy as np
2 from sklearn.neighbors import KNeighborsClassifier
3 from sklearn.model_selection import train_test_split
4 from sklearn.metrics import accuracy_score
5 from sklearn.cluster import KMeans
6 from sklearn.datasets import load_iris
7 iris = load_iris()
8 X = iris.data
9 y = iris.target
10
11 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
12
```

```
In [44]: 1 k = 3 # Set the value of k
2 knn_classifier = KNeighborsClassifier(n_neighbors=k)
3 knn_classifier.fit(X_train, y_train)
```

Out[44]: KNeighborsClassifier(n_neighbors=3)

```
In [48]: 1 y_pred = knn_classifier.predict(X_test)
          2 y_pred
          3
```

C:\Users\Anusha V\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

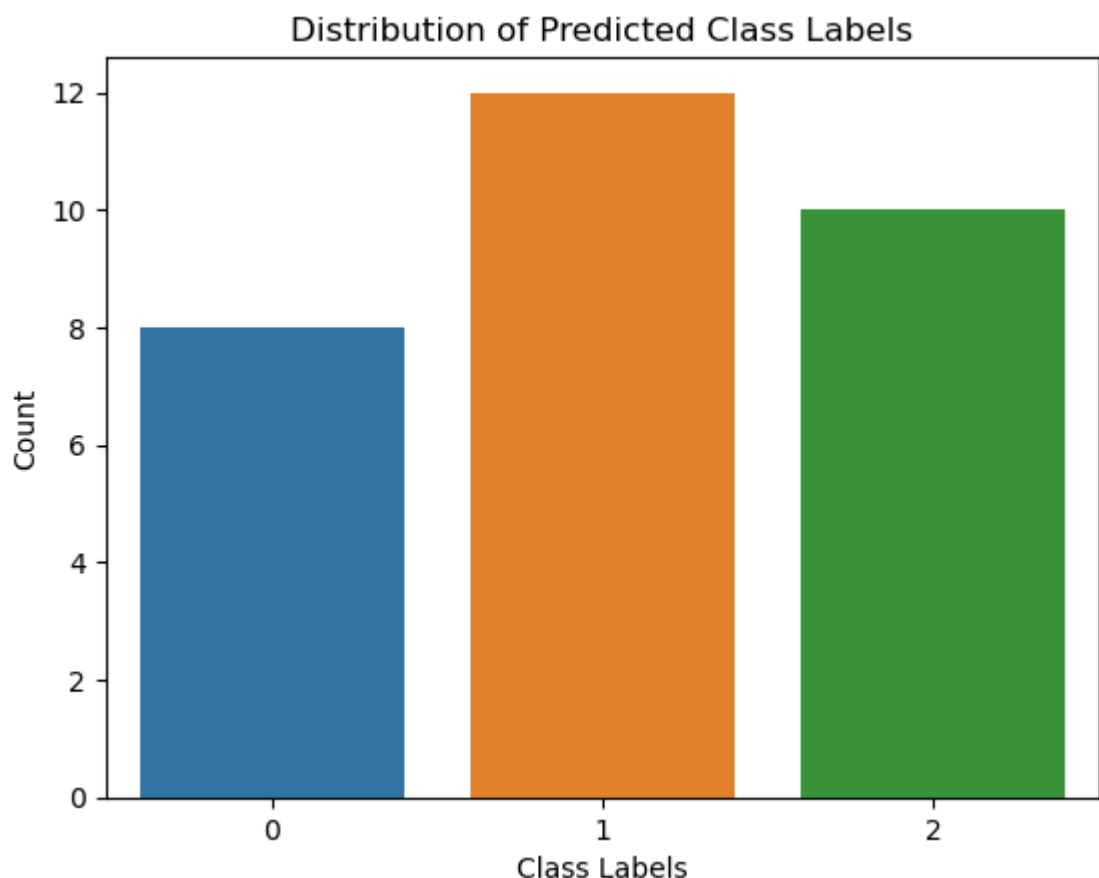
```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
Out[48]: array([0, 1, 2, 2, 1, 2, 1, 1, 1, 0, 1, 0, 0, 2, 1, 2, 2, 2, 1, 1, 2, 2,
                1, 0, 1, 0, 0, 2, 0, 1])
```


```
In [49]: 1 import seaborn as sns
          2 import matplotlib.pyplot as plt
          3 sns.countplot(y_pred)
          4 plt.title("Distribution of Predicted Class Labels")
          5 plt.xlabel("Class Labels")
          6 plt.ylabel("Count")
          7 plt.show()
```

C:\Users\Anusha V\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```



```
In [50]: 1 from sklearn.cluster import KMeans
2 import matplotlib.pyplot as plt
3
4 # Generate some example data (replace this with your own data)
5 import numpy as np
6 X = np.array([[1, 2], [5, 8], [1.5, 1.8], [8, 8], [1, 0.6], [9, 11]])
7
8 # Create a K-Means model with the desired number of clusters (e.g., 2)
9 kmeans = KMeans(n_clusters=2)
10
11 # Fit the model to your data
12 kmeans.fit(X)
13
14 # Get the cluster labels for each data point
15 labels = kmeans.labels_
16
17 # Get the coordinates of the cluster centers
18 centers = kmeans.cluster_centers_
19
20 # Plot the data points and cluster centers
21 plt.scatter(X[:, 0], X[:, 1], c=labels)
22 plt.scatter(centers[:, 0], centers[:, 1], marker='x', s=200, linewidths=2)
23 plt.show()
```




Simple linear regression

```

In [7]: 1 import pandas as pd
        2 import numpy as np
        3 from sklearn.model_selection import train_test_split
        4 from sklearn.linear_model import LinearRegression
        5 import matplotlib.pyplot as plt
        6 data1= pd.read_csv(r"C:\Users\Anusha V\Downloads\fish.csv")
        7
        8 X = data1[['nofish']]
        9 y = data1['persons']
       10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2
       11 model = LinearRegression()
       12 model.fit(X_train, y_train)
       13 y_pred = model.predict(X_test)
       14 plt.scatter(X, y, color='blue', label='Data')
       15 plt.plot(X_test, y_pred, color='red', linewidth=2, label='Regression Li
       16 plt.xlabel('Independent Variable')
       17 plt.ylabel('Dependent Variable')
       18 plt.legend()
       19 plt.show()
       20 slope = model.coef_[0]
       21 intercept = model.intercept_
       22 print(f"Slope (m): {slope}")
       23 print(f"Intercept (b): {intercept}")
       24

```



Slope (m): -0.07705253035220574
Intercept (b): 2.517730496453901

Breast Cancer Wisconsin (diagnostic)¶

```
In [31]: 1 import numpy as np
          2 import pandas as pd
          3 from sklearn.datasets import load_breast_cancer
          4 from sklearn.model_selection import train_test_split
          5 from sklearn.ensemble import RandomForestClassifier
          6 from sklearn.metrics import accuracy_score, classification_report
```

```
In [32]: 1 data = load_breast_cancer()
          2 X = data.data
          3 y = data.target
```

```
In [33]: 1 rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
          2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [34]: 1 rf_classifier.fit(X_train, y_train)
          2
```

Out[34]: RandomForestClassifier(random_state=42)

```
In [35]: 1 y_pred = rf_classifier.predict(X_test)
          2 accuracy = accuracy_score(y_test, y_pred)
          3 accuracy
```

Out[35]: 0.9649122807017544

```
In [36]: 1 report = classification_report(y_test, y_pred)
          2 print(report)
```

	precision	recall	f1-score	support
0	0.98	0.93	0.95	43
1	0.96	0.99	0.97	71
accuracy			0.96	114
macro avg	0.97	0.96	0.96	114
weighted avg	0.97	0.96	0.96	114

Build decision tree-based model in python

like Breast Cancer Wisconsin (diagnostic) dataset from sci-kit learn Or any classification dataset from UCI , Kaggle

```
In [40]: 1 from sklearn.datasets import load_breast_cancer
2 from sklearn.model_selection import train_test_split
3 from sklearn.tree import DecisionTreeClassifier
4 from sklearn.metrics import accuracy_score, classification_report
5 breast_cancer = load_breast_cancer()
6 X = breast_cancer.data # Features
7 y = breast_cancer.target # Target variable
8 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
9 clf = DecisionTreeClassifier(random_state=42)
10 clf.fit(X_train, y_train)
11 y_pred = clf.predict(X_test)
12 accuracy = accuracy_score(y_test, y_pred)
13 print(f"Accuracy: {accuracy:.2f}")
14 print("Classification Report:")
15 print(classification_report(y_test, y_pred))
16
17
```

Accuracy: 0.95

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.93	0.93	43
1	0.96	0.96	0.96	71
accuracy			0.95	114
macro avg	0.94	0.94	0.94	114
weighted avg	0.95	0.95	0.95	114

Build Random Forest-based model in python for

Breast Cancer Wisconsin (diagnostic) dataset from sci-kit learn Or dataset from UCI , Kaggle

```
In [41]: 1 from sklearn.datasets import load_breast_cancer
2 from sklearn.model_selection import train_test_split
3 from sklearn.ensemble import RandomForestClassifier
4 from sklearn.metrics import accuracy_score, classification_report
5 breast_cancer = load_breast_cancer()
6 X = breast_cancer.data # Features
7 y = breast_cancer.target # Target variable
8 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2
9 rf_classifier = RandomForestClassifier(n_estimators=100, random_state=4
10 rf_classifier.fit(X_train, y_train)
11 y_pred = rf_classifier.predict(X_test)
12 accuracy = accuracy_score(y_test, y_pred)
13 print(f"Accuracy: {accuracy:.2f}")
14 print("Classification Report:")
15 print(classification_report(y_test, y_pred))
16
```

Accuracy: 0.96

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.93	0.95	43
1	0.96	0.99	0.97	71
accuracy			0.96	114
macro avg	0.97	0.96	0.96	114
weighted avg	0.97	0.96	0.96	114

Iris dataset from sci-kit learn Perform data exploration,

preprocessing and splitting

```
In [29]: 1 from sklearn.datasets import load_iris
2 from sklearn.model_selection import train_test_split
3 import pandas as pd
4 iris = load_iris()
5 iris_df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
6 iris_df['target'] = iris.target
7 print(iris_df.head())
8 X = iris.data # Features
9 y = iris.target # Target variable
10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
11 print("X_train shape:", X_train.shape)
12 print("X_test shape:", X_test.shape)
13 print("y_train shape:", y_train.shape)
14 print("y_test shape:", y_test.shape)
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
0	5.1	3.5	1.4	0.
1	4.9	3.0	1.4	0.
2	4.7	3.2	1.3	0.
3	4.6	3.1	1.5	0.
4	5.0	3.6	1.4	0.

	target
0	0
1	0
2	0
3	0
4	0

X_train shape: (120, 4)
X_test shape: (30, 4)
y_train shape: (120,)
y_test shape: (30,)

Regression Model with Deep Learning:


```

In [22]: 1 import numpy as np
          2 from sklearn.datasets import load_boston
          3 from sklearn.model_selection import train_test_split
          4 from sklearn.preprocessing import StandardScaler
          5 from tensorflow.keras import models, layers
          6
          7 # Load the Boston Housing Prices dataset
          8 boston = load_boston()
          9 data = boston.data
         10 targets = boston.target
         11
         12 # Split the data into training and testing sets
         13 X_train, X_test, y_train, y_test = train_test_split(data, targets, test
         14
         15 # Standardize the data
         16 scaler = StandardScaler()
         17 X_train = scaler.fit_transform(X_train)
         18 X_test = scaler.transform(X_test)
         19
         20 # Build the regression model
         21 model_reg = models.Sequential()
         22 model_reg.add(layers.Dense(64, activation='relu', input_shape=(X_train.
         23 model_reg.add(layers.Dense(1)) # Output layer for regression
         24
         25 # Compile the model
         26 model_reg.compile(optimizer='adam', loss='mean_squared_error', metrics=
         27
         28 # Train the model
         29 model_reg.fit(X_train, y_train, epochs=50, batch_size=16, validation_da
         30
26/26 [=====] - 0s 19ms/step - loss: 27.8384 -
mae: 3.8863 - val_loss: 29.0147 - val_mae: 3.6542
Epoch 27/50
26/26 [=====] - 0s 14ms/step - loss: 27.1955 -
mae: 3.8508 - val_loss: 28.2891 - val_mae: 3.6082
Epoch 28/50
26/26 [=====] - 0s 13ms/step - loss: 26.5516 -
mae: 3.8142 - val_loss: 27.7366 - val_mae: 3.5744
Epoch 29/50
26/26 [=====] - 0s 10ms/step - loss: 26.0652 -
mae: 3.7790 - val_loss: 27.1639 - val_mae: 3.5189
Epoch 30/50
26/26 [=====] - 0s 11ms/step - loss: 25.4592 -
mae: 3.7307 - val_loss: 26.7644 - val_mae: 3.4852
Epoch 31/50
26/26 [=====] - 0s 11ms/step - loss: 24.9334 -
mae: 3.7008 - val_loss: 26.1122 - val_mae: 3.4524
Epoch 32/50
26/26 [=====] - 0s 9ms/step - loss: 24.4530 - m
ae: 3.6647 - val_loss: 25.7499 - val_mae: 3.4131

```

Classification Model with Deep Learning:

```
In [23]: 1 from sklearn.datasets import load_iris
2 from sklearn.model_selection import train_test_split
3 from sklearn.preprocessing import StandardScaler, LabelEncoder
4 from tensorflow.keras import models, layers, utils
5
6 # Load the Iris dataset
7 iris = load_iris()
8 data = iris.data
9 targets = iris.target
10
11 # Split the data into training and testing sets
12 X_train, X_test, y_train, y_test = train_test_split(data, targets, test
13
14 # Standardize the data
15 scaler = StandardScaler()
16 X_train = scaler.fit_transform(X_train)
17 X_test = scaler.transform(X_test)
18
19 # One-hot encode the labels
20 y_train = utils.to_categorical(y_train)
21 y_test = utils.to_categorical(y_test)
22
23 # Build the classification model
24 model_cls = models.Sequential()
25 model_cls.add(layers.Dense(64, activation='relu', input_shape=(X_train.
26 model_cls.add(layers.Dense(3, activation='softmax'))) # Output Layer fo
27
28 # Compile the model
29 model_cls.compile(optimizer='adam', loss='categorical_crossentropy', me
30
31 # Train the model
32 model_cls.fit(X_train, y_train, epochs=50, batch_size=16, validation_da
33
```

```
Epoch 1/50
8/8 [=====] - 2s 77ms/step - loss: 1.0962 - accuracy: 0.5750 - val_loss: 1.0760 - val_accuracy: 0.6000
Epoch 2/50
8/8 [=====] - 0s 18ms/step - loss: 0.9886 - accuracy: 0.6250 - val_loss: 0.9611 - val_accuracy: 0.6000
Epoch 3/50
8/8 [=====] - 0s 27ms/step - loss: 0.8924 - accuracy: 0.6333 - val_loss: 0.8628 - val_accuracy: 0.6000
Epoch 4/50
8/8 [=====] - 0s 28ms/step - loss: 0.8104 - accuracy: 0.6583 - val_loss: 0.7776 - val_accuracy: 0.8000
Epoch 5/50
8/8 [=====] - 0s 19ms/step - loss: 0.7405 - accuracy: 0.7583 - val_loss: 0.7028 - val_accuracy: 0.8333
Epoch 6/50
8/8 [=====] - 0s 18ms/step - loss: 0.6797 - accuracy: 0.8000 - val_loss: 0.6372 - val_accuracy: 0.8667
Epoch 7/50
8/8 [=====] - 0s 21ms/step - loss: 0.6272 - accuracy: 0.8167 - val_loss: 0.5803 - val_accuracy: 0.9000
Epoch 8/50
8/8 [=====] - 0s 21ms/step - loss: 0.5848 - accuracy: 0.8083 - val_loss: 0.5305 - val_accuracy: 0.9000
Epoch 9/50
8/8 [=====] - 0s 26ms/step - loss: 0.5440 - accuracy: 0.8083 - val_loss: 0.4928 - val_accuracy: 0.9000
Epoch 10/50
8/8 [=====] - 0s 25ms/step - loss: 0.5150 - accuracy: 0.8167 - val_loss: 0.4574 - val_accuracy: 0.9000
Epoch 11/50
8/8 [=====] - 0s 25ms/step - loss: 0.4879 - accuracy: 0.8167 - val_loss: 0.4292 - val_accuracy: 0.9000
Epoch 12/50
8/8 [=====] - 0s 25ms/step - loss: 0.4653 - accuracy: 0.8167 - val_loss: 0.4060 - val_accuracy: 0.9000
Epoch 13/50
8/8 [=====] - 0s 26ms/step - loss: 0.4465 - accuracy: 0.8167 - val_loss: 0.3853 - val_accuracy: 0.9000
Epoch 14/50
8/8 [=====] - 0s 12ms/step - loss: 0.4301 - accuracy: 0.8167 - val_loss: 0.3672 - val_accuracy: 0.9000
Epoch 15/50
8/8 [=====] - 0s 12ms/step - loss: 0.4154 - accuracy: 0.8167 - val_loss: 0.3516 - val_accuracy: 0.9000
Epoch 16/50
8/8 [=====] - 0s 11ms/step - loss: 0.4025 - accuracy: 0.8250 - val_loss: 0.3384 - val_accuracy: 0.9000
Epoch 17/50
8/8 [=====] - 0s 11ms/step - loss: 0.3915 - accuracy: 0.8250 - val_loss: 0.3253 - val_accuracy: 0.9000
Epoch 18/50
8/8 [=====] - 0s 22ms/step - loss: 0.3807 - accuracy: 0.8250 - val_loss: 0.3140 - val_accuracy: 0.9000
Epoch 19/50
8/8 [=====] - 0s 13ms/step - loss: 0.3711 - accuracy: 0.8333 - val_loss: 0.3046 - val_accuracy: 0.9000
Epoch 20/50
8/8 [=====] - 0s 16ms/step - loss: 0.3625 - accuracy: 0.8333 - val_loss: 0.2951 - val_accuracy: 0.9000
Epoch 21/50
```

```
8/8 [=====] - 0s 31ms/step - loss: 0.3544 - accur
acy: 0.8333 - val_loss: 0.2864 - val_accuracy: 0.9000
Epoch 22/50
8/8 [=====] - 0s 16ms/step - loss: 0.3466 - accur
acy: 0.8417 - val_loss: 0.2783 - val_accuracy: 0.9000
Epoch 23/50
8/8 [=====] - 0s 13ms/step - loss: 0.3391 - accur
acy: 0.8417 - val_loss: 0.2716 - val_accuracy: 0.9000
Epoch 24/50
8/8 [=====] - 0s 15ms/step - loss: 0.3318 - accur
acy: 0.8417 - val_loss: 0.2645 - val_accuracy: 0.9000
Epoch 25/50
8/8 [=====] - 0s 14ms/step - loss: 0.3252 - accur
acy: 0.8417 - val_loss: 0.2577 - val_accuracy: 0.9000
Epoch 26/50
8/8 [=====] - 0s 13ms/step - loss: 0.3189 - accur
acy: 0.8583 - val_loss: 0.2509 - val_accuracy: 0.9000
Epoch 27/50
8/8 [=====] - 0s 13ms/step - loss: 0.3122 - accur
acy: 0.8583 - val_loss: 0.2441 - val_accuracy: 0.9000
Epoch 28/50
8/8 [=====] - 0s 12ms/step - loss: 0.3060 - accur
acy: 0.8667 - val_loss: 0.2387 - val_accuracy: 0.9000
Epoch 29/50
8/8 [=====] - 0s 11ms/step - loss: 0.3000 - accur
acy: 0.8667 - val_loss: 0.2339 - val_accuracy: 0.9000
Epoch 30/50
8/8 [=====] - 0s 19ms/step - loss: 0.2946 - accur
acy: 0.8667 - val_loss: 0.2284 - val_accuracy: 0.9000
Epoch 31/50
8/8 [=====] - 0s 13ms/step - loss: 0.2890 - accur
acy: 0.8667 - val_loss: 0.2237 - val_accuracy: 0.9000
Epoch 32/50
8/8 [=====] - 0s 26ms/step - loss: 0.2833 - accur
acy: 0.8667 - val_loss: 0.2187 - val_accuracy: 0.9000
Epoch 33/50
8/8 [=====] - 0s 17ms/step - loss: 0.2778 - accur
acy: 0.8833 - val_loss: 0.2142 - val_accuracy: 0.9000
Epoch 34/50
8/8 [=====] - 0s 17ms/step - loss: 0.2731 - accur
acy: 0.8917 - val_loss: 0.2102 - val_accuracy: 0.9667
Epoch 35/50
8/8 [=====] - 0s 24ms/step - loss: 0.2670 - accur
acy: 0.9000 - val_loss: 0.2049 - val_accuracy: 0.9667
Epoch 36/50
8/8 [=====] - 0s 18ms/step - loss: 0.2614 - accur
acy: 0.9000 - val_loss: 0.2012 - val_accuracy: 0.9667
Epoch 37/50
8/8 [=====] - 0s 13ms/step - loss: 0.2564 - accur
acy: 0.9000 - val_loss: 0.1968 - val_accuracy: 0.9667
Epoch 38/50
8/8 [=====] - 0s 16ms/step - loss: 0.2518 - accur
acy: 0.9083 - val_loss: 0.1921 - val_accuracy: 0.9667
Epoch 39/50
8/8 [=====] - 0s 13ms/step - loss: 0.2465 - accur
acy: 0.9083 - val_loss: 0.1880 - val_accuracy: 0.9667
Epoch 40/50
8/8 [=====] - 0s 16ms/step - loss: 0.2415 - accur
acy: 0.9083 - val_loss: 0.1835 - val_accuracy: 0.9667
Epoch 41/50
8/8 [=====] - 0s 15ms/step - loss: 0.2376 - accur
```

```

acy: 0.9167 - val_loss: 0.1780 - val_accuracy: 0.9667
Epoch 42/50
8/8 [=====] - 0s 18ms/step - loss: 0.2329 - accur
acy: 0.9167 - val_loss: 0.1739 - val_accuracy: 0.9667
Epoch 43/50
8/8 [=====] - 0s 16ms/step - loss: 0.2272 - accur
acy: 0.9250 - val_loss: 0.1711 - val_accuracy: 0.9667
Epoch 44/50
8/8 [=====] - 0s 13ms/step - loss: 0.2237 - accur
acy: 0.9250 - val_loss: 0.1689 - val_accuracy: 0.9667
Epoch 45/50
8/8 [=====] - 0s 13ms/step - loss: 0.2183 - accur
acy: 0.9250 - val_loss: 0.1641 - val_accuracy: 0.9667
Epoch 46/50
8/8 [=====] - 0s 13ms/step - loss: 0.2138 - accur
acy: 0.9333 - val_loss: 0.1597 - val_accuracy: 0.9667
Epoch 47/50
8/8 [=====] - 0s 20ms/step - loss: 0.2097 - accur
acy: 0.9417 - val_loss: 0.1568 - val_accuracy: 0.9667
Epoch 48/50
8/8 [=====] - 0s 33ms/step - loss: 0.2052 - accur
acy: 0.9417 - val_loss: 0.1526 - val_accuracy: 0.9667
Epoch 49/50
8/8 [=====] - 0s 20ms/step - loss: 0.2008 - accur
acy: 0.9417 - val_loss: 0.1496 - val_accuracy: 0.9667
Epoch 50/50
8/8 [=====] - 0s 26ms/step - loss: 0.1971 - accur
acy: 0.9500 - val_loss: 0.1473 - val_accuracy: 0.9667

```

Out[23]: <keras.src.callbacks.History at 0x1d313fb0a60>

Analyzing Performance:

```

In [26]: 1 # Build the regression model
          2 model_reg = models.Sequential()
          3 model_reg.add(layers.Dense(64, activation='relu', input_shape=(X_train.
          4 model_reg.add(layers.Dense(1)) # Output layer for regression
          5
          6 # Compile the model
          7 model_reg.compile(optimizer='adam', loss='mean_squared_error', metrics=
          8 model_reg.compile

```

Out[26]: <bound method Model.compile of <keras.src.engine.sequential.Sequential object at 0x000001D320980490>>

Create a model to analyses the relation between CIE and SEE result

In [27]:

```
1 import numpy as np
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from sklearn.preprocessing import StandardScaler
5 from tensorflow.keras import models, layers
6
7 # Assuming you have a dataset with columns 'CIE' and 'SEE'
8 # Load your dataset
9 # For example, assuming you have a CSV file named 'your_dataset.csv'
10 # df = pd.read_csv('your_dataset.csv')
11
12 # For the purpose of this example, let's generate some random data
13 np.random.seed(42)
14 num_samples = 1000
15 cie_data = np.random.rand(num_samples) * 10 # Random CIE values between
16 see_data = 2 * cie_data + np.random.randn(num_samples) * 2 # Linear re
17
18 # Create a DataFrame
19 df = pd.DataFrame({'CIE': cie_data, 'SEE': see_data})
20
21 # Split the data into training and testing sets
22 train_data, test_data = train_test_split(df, test_size=0.2, random_stat
23
24 # Standardize the data
25 scaler = StandardScaler()
26 X_train = scaler.fit_transform(train_data[['CIE']])
27 y_train = scaler.fit_transform(train_data[['SEE']])
28 X_test = scaler.transform(test_data[['CIE']])
29 y_test = scaler.transform(test_data[['SEE']])
30
31 # Build the regression model
32 model = models.Sequential()
33 model.add(layers.Dense(64, activation='relu', input_shape=(1,)))
34 model.add(layers.Dense(1)) # Output layer for regression
35
36 # Compile the model
37 model.compile(optimizer='adam', loss='mean_squared_error', metrics=['ma
38
39 # Train the model
40 model.fit(X_train, y_train, epochs=50, batch_size=16, validation_data=(
41
```

50/50 [=====] - 0s 7ms/step - loss: 0.1081 - ma

Create a model to analyze the relation between crop yield and rain fall rate

In [28]:

```
1 import numpy as np
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from sklearn.preprocessing import StandardScaler
5 from tensorflow.keras import models, layers
6
7 # Assuming you have a dataset with columns 'Crop_Yield' and 'Rainfall_R
8 # Load your dataset
9 # For example, assuming you have a CSV file named 'your_dataset.csv'
10 # df = pd.read_csv('your_dataset.csv')
11
12 # For the purpose of this example, let's generate some random data
13 np.random.seed(42)
14 num_samples = 1000
15 rainfall_data = np.random.rand(num_samples) * 100 # Random rainfall va
16 crop_yield_data = 2 * rainfall_data + np.random.randn(num_samples) * 10
17
18 # Create a DataFrame
19 df = pd.DataFrame({'Rainfall_Rate': rainfall_data, 'Crop_Yield': crop_y
20
21 # Split the data into training and testing sets
22 train_data, test_data = train_test_split(df, test_size=0.2, random_stat
23
24 # Standardize the data
25 scaler = StandardScaler()
26 X_train = scaler.fit_transform(train_data[['Rainfall_Rate']])
27 y_train = scaler.fit_transform(train_data[['Crop_Yield']])
28 X_test = scaler.transform(test_data[['Rainfall_Rate']])
29 y_test = scaler.transform(test_data[['Crop_Yield']])
30
31 # Build the regression model
32 model = models.Sequential()
33 model.add(layers.Dense(64, activation='relu', input_shape=(1,)))
34 model.add(layers.Dense(1)) # Output layer for regression
35
36 # Compile the model
37 model.compile(optimizer='adam', loss='mean_squared_error', metrics=['ma
38
39 # Train the model
40 model.fit(X_train, y_train, epochs=50, batch_size=16, validation_data=(
41
```


e: 0.1466 - val loss: 0.9855 - val mae: 0.8362

Build linear regression model using

- Stats model
- Scikit learn

Using Statsmodels

In [29]:

```
1 import statsmodels.api as sm
2 import pandas as pd
3 import numpy as np
4 from sklearn.model_selection import train_test_split
5 from sklearn.preprocessing import StandardScaler
6
7 # Assuming you have a dataset with columns 'Rainfall_Rate' and 'Crop_Yi
8 # Load your dataset
9 # For example, assuming you have a CSV file named 'your_dataset.csv'
10 # df = pd.read_csv('your_dataset.csv')
11
12 # For the purpose of this example, let's generate some random data
13 np.random.seed(42)
14 num_samples = 1000
15 rainfall_data = np.random.rand(num_samples) * 100 # Random rainfall va
16 crop_yield_data = 2 * rainfall_data + np.random.randn(num_samples) * 10
17
18 # Create a DataFrame
19 df = pd.DataFrame({'Rainfall_Rate': rainfall_data, 'Crop_Yield': crop_y
20
21 # Split the data into training and testing sets
22 train_data, test_data = train_test_split(df, test_size=0.2, random_stat
23
24 # Standardize the data
25 scaler = StandardScaler()
26 X_train = scaler.fit_transform(train_data[['Rainfall_Rate']])
27 y_train = train_data['Crop_Yield']
28
29 # Add a constant term for the intercept
30 X_train = sm.add_constant(X_train)
31
32 # Create and fit the model
33 model = sm.OLS(y_train, X_train)
34 result = model.fit()
35
36 # Display the summary of the regression
37 print(result.summary())
38
```

```

=====
OLS Regression Results
=====
====
Dep. Variable:          Crop_Yield    R-squared:
0.971
Model:                  OLS          Adj. R-squared:
0.971
Method:                 Least Squares    F-statistic:          2.673
e+04
Date:                  Sun, 26 Nov 2023    Prob (F-statistic):
0.00
Time:                  13:54:17    Log-Likelihood:          -29
79.0
No. Observations:      800    AIC:          5
962.
Df Residuals:          798    BIC:          5
971.
Df Model:              1
Covariance Type:       nonrobust
=====
====
               coef    std err          t      P>|t|      [0.025    0.
975]
-----
----
const          100.0905      0.355    282.119      0.000      99.394    10
0.787
x1              58.0053      0.355    163.496      0.000      57.309      5
8.702
=====
====
Omnibus:              1.240    Durbin-Watson:
1.924
Prob(Omnibus):        0.538    Jarque-Bera (JB):
1.235
Skew:                 0.014    Prob(JB):
0.539
Kurtosis:             2.810    Cond. No.
1.00
=====
====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is cor
rectly specified.
```

Using Scikit-learn:

In [30]:

```
1 from sklearn.linear_model import LinearRegression
2 from sklearn.model_selection import train_test_split
3 from sklearn.preprocessing import StandardScaler
4 import numpy as np
5
6 # Assuming you have a dataset with columns 'Rainfall_Rate' and 'Crop_Yield'
7 # Load your dataset
8 # For example, assuming you have a CSV file named 'your_dataset.csv'
9 # df = pd.read_csv('your_dataset.csv')
10
11 # For the purpose of this example, let's generate some random data
12 np.random.seed(42)
13 num_samples = 1000
14 rainfall_data = np.random.rand(num_samples) * 100 # Random rainfall values
15 crop_yield_data = 2 * rainfall_data + np.random.randn(num_samples) * 10
16
17 # Create a DataFrame
18 df = pd.DataFrame({'Rainfall_Rate': rainfall_data, 'Crop_Yield': crop_yield_data})
19
20 # Split the data into training and testing sets
21 X_train, X_test, y_train, y_test = train_test_split(df[['Rainfall_Rate']], df['Crop_Yield'],
22                                                    test_size=0.2, random_state=42)
23
24 # Standardize the data (not strictly necessary for Linear regression, but helps)
25 scaler = StandardScaler()
26 X_train_scaled = scaler.fit_transform(X_train)
27 X_test_scaled = scaler.transform(X_test)
28
29 # Create and fit the model
30 model = LinearRegression()
31 model.fit(X_train_scaled, y_train)
32
33 # Display the coefficients and intercept
34 print("Coefficients:", model.coef_)
35 print("Intercept:", model.intercept_)
```

Coefficients: [58.00531307]

Intercept: 100.09054272144182

Implementation in python

- Build regression model
- Evaluate the model
- To minimize the cost function


```

In [31]: 1 import numpy as np
          2 import matplotlib.pyplot as plt
          3 from sklearn.model_selection import train_test_split
          4 from sklearn.metrics import mean_squared_error
          5
          6 # Generate some random data for demonstration
          7 np.random.seed(42)
          8 X = 2 * np.random.rand(100, 1)
          9 y = 4 + 3 * X + np.random.randn(100, 1)
         10
         11 # Split the data into training and testing sets
         12 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
         13
         14 # Add a bias term to the features
         15 X_train_bias = np.c_[np.ones((len(X_train), 1)), X_train]
         16 X_test_bias = np.c_[np.ones((len(X_test), 1)), X_test]
         17
         18 # Initialize random coefficients
         19 theta = np.random.randn(2, 1)
         20
         21 # Define the learning rate and number of iterations
         22 learning_rate = 0.01
         23 n_iterations = 1000
         24
         25 # Implement gradient descent to minimize the cost function
         26 for iteration in range(n_iterations):
         27     gradients = 2/len(X_train) * X_train_bias.T.dot(X_train_bias.dot(theta) - y_train)
         28     theta = theta - learning_rate * gradients
         29
         30 # Evaluate the model on the test set
         31 y_pred = X_test_bias.dot(theta)
         32
         33 # Calculate the mean squared error
         34 mse = mean_squared_error(y_test, y_pred)
         35
         36 # Print the Learned coefficients and MSE
         37 print("Learned Coefficients:", theta.flatten())
         38 print("Mean Squared Error on Test Set:", mse)
         39
         40 # Visualize the data and the Learned regression line
         41 plt.scatter(X_test, y_test, label='Test Data')
         42 plt.plot(X_test, y_pred, color='red', label='Regression Line')
         43 plt.xlabel('X')
         44 plt.ylabel('y')
         45 plt.legend()
         46 plt.show()
         47

```

Learned Coefficients: [4.08358595 2.85220796]

Mean Squared Error on Test Set: 0.6645849137924722



**build Logistic regression model in
pythonEvaluation and optimization of the
model**

In [32]:

```
1 import numpy as np
2 import pandas as pd
3 from sklearn.model_selection import train_test_split, GridSearchCV
4 from sklearn.preprocessing import StandardScaler
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
7
8 # Assuming you have a dataset with features and labels
9 # For example, assuming you have a CSV file named 'your_dataset.csv'
10 # df = pd.read_csv('your_dataset.csv')
11
12 # For the purpose of this example, let's generate some random data
13 np.random.seed(42)
14 num_samples = 1000
15 X = np.random.rand(num_samples, 2) * 10 # Random features
16 y = (X[:, 0] + X[:, 1] > 10).astype(int) # Binary classification task
17
18 # Split the data into training and testing sets
19 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
20
21 # Standardize the data
22 scaler = StandardScaler()
23 X_train_scaled = scaler.fit_transform(X_train)
24 X_test_scaled = scaler.transform(X_test)
25
26 # Build the logistic regression model
27 model = LogisticRegression()
28
29 # Train the model
30 model.fit(X_train_scaled, y_train)
31
32 # Evaluate the model on the test set
33 y_pred = model.predict(X_test_scaled)
34
35 # Calculate accuracy
36 accuracy = accuracy_score(y_test, y_pred)
37
38 # Print confusion matrix and classification report
39 conf_matrix = confusion_matrix(y_test, y_pred)
40 class_report = classification_report(y_test, y_pred)
41
42 print("Accuracy:", accuracy)
43 print("Confusion Matrix:\n", conf_matrix)
44 print("Classification Report:\n", class_report)
45
46 # Hyperparameter tuning using GridSearchCV
47 param_grid = {'C': [0.001, 0.01, 0.1, 1, 10, 100]} # Regularization parameter
48 grid_search = GridSearchCV(LogisticRegression(), param_grid, cv=5)
49 grid_search.fit(X_train_scaled, y_train)
50
51 # Print the best hyperparameters
52 print("Best Hyperparameters:", grid_search.best_params_)
53
54 # Evaluate the model with the best hyperparameters
55 best_model = grid_search.best_estimator_
56 y_pred_best = best_model.predict(X_test_scaled)
57
58 # Calculate accuracy for the optimized model
59 accuracy_best = accuracy_score(y_test, y_pred_best)
60
61 # Print confusion matrix and classification report for the optimized model
```



```

62 conf_matrix_best = confusion_matrix(y_test, y_pred_best)
63 class_report_best = classification_report(y_test, y_pred_best)
64
65 print("Optimized Model Accuracy:", accuracy_best)
66 print("Optimized Model Confusion Matrix:\n", conf_matrix_best)
67 print("Optimized Model Classification Report:\n", class_report_best)
68

```

Accuracy: 1.0

Confusion Matrix:

```

[[106  0]
 [  0  94]]

```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	106
1	1.00	1.00	1.00	94
accuracy			1.00	200
macro avg	1.00	1.00	1.00	200
weighted avg	1.00	1.00	1.00	200

Best Hyperparameters: {'C': 1}

Optimized Model Accuracy: 1.0

Optimized Model Confusion Matrix:

```

[[106  0]
 [  0  94]]

```

Optimized Model Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	106
1	1.00	1.00	1.00	94
accuracy			1.00	200
macro avg	1.00	1.00	1.00	200
weighted avg	1.00	1.00	1.00	200

Feature scaling with StandardScaler() or other method Dropping unnecessary features

Data splitting Dealing with imbalanced dataset

```
In [34]: 1 pip install pandas scikit-learn imbalanced-learn
          2
```

0:20:14

WARNING: Retrying (Retry(total=4, connect=None, read=None, redirect=None, status=None)) after connection broken by 'ReadTimeoutError("HTTPSConnectionPool(host='pypi.org', port=443): Read timed out. (read timeout=15)")': /simple/numpy/

WARNING: Retrying (Retry(total=3, connect=None, read=None, redirect=None, status=None)) after connection broken by 'ReadTimeoutError("HTTPSConnectionPool(host='pypi.org', port=443): Read timed out. (read timeout=15)")': /simple/numpy/

WARNING: Retrying (Retry(total=2, connect=None, read=None, redirect=None, status=None)) after connection broken by 'ReadTimeoutError("HTTPSConnectionPool(host='pypi.org', port=443): Read timed out. (read timeout=15)")': /simple/numpy/

WARNING: Retrying (Retry(total=1, connect=None, read=None, redirect=None, status=None)) after connection broken by 'ReadTimeoutError("HTTPSConnectionPool(host='pypi.org', port=443): Read timed out. (read timeout=15)")': /simple/numpy/

WARNING: Retrying (Retry(total=0, connect=None, read=None, redirect=None, status=None)) after connection broken by 'ReadTimeoutError("HTTPSConn

```

In [36]: 1 import pandas as pd
2 from sklearn.model_selection import train_test_split
3 from sklearn.preprocessing import StandardScaler
4 from imblearn.over_sampling import SMOTE
5 from sklearn.datasets import make_classification
6
7 # Generate synthetic imbalanced data for demonstration
8 X, y = make_classification(n_samples=1000, n_features=10, n_classes=2,
9
10 # Create a DataFrame (replace this with your actual dataset loading logic)
11 df = pd.DataFrame(X, columns=[f'feature_{i}' for i in range(X.shape[1])])
12 df['target'] = y
13
14 # Check existing column names
15 print("Existing Columns:", df.columns)
16
17 # Drop unnecessary features (replace 'feature_to_drop' with actual feature names)
18 features_to_drop = ['feature_to_drop']
19
20 # Check if the columns to drop exist in the DataFrame
21 columns_to_drop = [col for col in features_to_drop if col in df.columns]
22 df = df.drop(columns=columns_to_drop)
23
24 # Separate features (X) and target variable (y)
25 X = df.drop(columns=['target'])
26 y = df['target']
27
28 # Feature scaling using StandardScaler
29 scaler = StandardScaler()
30 X_scaled = scaler.fit_transform(X)
31
32 # Data splitting into training and testing sets
33 X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
34
35 # Dealing with an imbalanced dataset using SMOTE
36 smote = SMOTE(random_state=42)
37 X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)
38
39 # Now you can use X_train_resampled and y_train_resampled for training
40 # Remember to use X_test for evaluating the model, not X_train
41

```

```

Existing Columns: Index(['feature_0', 'feature_1', 'feature_2', 'feature_3', 'feature_4',
                        'feature_5', 'feature_6', 'feature_7', 'feature_8', 'feature_9',
                        'target'],
                        dtype='object')

```

Building Shallow Neural Network with Keras Dense Layer

In [37]:

```
1 import numpy as np
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from sklearn.preprocessing import StandardScaler
5 from sklearn.metrics import accuracy_score
6 from keras.models import Sequential
7 from keras.layers import Dense
8
9 # Assuming you have a dataset with features and labels
10 # For example, assuming you have a CSV file named 'your_dataset.csv'
11 # df = pd.read_csv('your_dataset.csv')
12
13 # For the purpose of this example, let's generate some random data
14 np.random.seed(42)
15 num_samples = 1000
16 X = np.random.rand(num_samples, 10) # Random features
17 y = np.random.randint(2, size=num_samples) # Binary classification labels
18
19 # Create a DataFrame
20 df = pd.DataFrame(X, columns=[f'feature_{i}' for i in range(X.shape[1])])
21 df['target'] = y
22
23 # Separate features (X) and target variable (y)
24 X = df.drop(columns=['target'])
25 y = df['target']
26
27 # Standardize the data
28 scaler = StandardScaler()
29 X_scaled = scaler.fit_transform(X)
30
31 # Split the data into training and testing sets
32 X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
33
34 # Build the shallow neural network model
35 model = Sequential()
36 model.add(Dense(units=16, input_dim=X_train.shape[1], activation='relu'))
37 model.add(Dense(units=1, activation='sigmoid'))
38
39 # Compile the model
40 model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
41
42 # Train the model
43 model.fit(X_train, y_train, epochs=10, batch_size=32, validation_split=0.2)
44
45 # Evaluate the model on the test set
46 y_pred_prob = model.predict(X_test)
47 y_pred = (y_pred_prob > 0.5).astype(int)
48
49 # Calculate accuracy
50 accuracy = accuracy_score(y_test, y_pred)
51 print("Test Accuracy:", accuracy)
52
```

```
Epoch 1/10
20/20 [=====] - 1s 17ms/step - loss: 0.7839 - accu
racy: 0.4828 - val_loss: 0.7774 - val_accuracy: 0.4875
Epoch 2/10
20/20 [=====] - 0s 7ms/step - loss: 0.7614 - accu
racy: 0.4844 - val_loss: 0.7619 - val_accuracy: 0.4688
Epoch 3/10
20/20 [=====] - 0s 7ms/step - loss: 0.7452 - accu
racy: 0.4891 - val_loss: 0.7473 - val_accuracy: 0.4688
Epoch 4/10
20/20 [=====] - 0s 6ms/step - loss: 0.7322 - accu
racy: 0.4891 - val_loss: 0.7357 - val_accuracy: 0.4688
Epoch 5/10
20/20 [=====] - 0s 7ms/step - loss: 0.7219 - accu
racy: 0.4953 - val_loss: 0.7263 - val_accuracy: 0.4750
Epoch 6/10
20/20 [=====] - 0s 6ms/step - loss: 0.7133 - accu
racy: 0.4953 - val_loss: 0.7188 - val_accuracy: 0.4688
Epoch 7/10
20/20 [=====] - 0s 7ms/step - loss: 0.7061 - accu
racy: 0.4984 - val_loss: 0.7139 - val_accuracy: 0.4625
Epoch 8/10
20/20 [=====] - 0s 5ms/step - loss: 0.7004 - accu
racy: 0.5063 - val_loss: 0.7085 - val_accuracy: 0.4938
Epoch 9/10
20/20 [=====] - 0s 7ms/step - loss: 0.6954 - accu
racy: 0.5312 - val_loss: 0.7047 - val_accuracy: 0.4938
Epoch 10/10
20/20 [=====] - 0s 6ms/step - loss: 0.6921 - accu
racy: 0.5375 - val_loss: 0.7025 - val_accuracy: 0.5000
7/7 [=====] - 0s 3ms/step
Test Accuracy: 0.505
```

Building Deep Neural Network with Keras Dense Layers

```
In [38]: 1 import numpy as np
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from sklearn.preprocessing import StandardScaler
5 from sklearn.metrics import accuracy_score
6 from keras.models import Sequential
7 from keras.layers import Dense
8
9 # Assuming you have a dataset with features and labels
10 # For example, assuming you have a CSV file named 'your_dataset.csv'
11 # df = pd.read_csv('your_dataset.csv')
12
13 # For the purpose of this example, let's generate some random data
14 np.random.seed(42)
15 num_samples = 1000
16 X = np.random.rand(num_samples, 10) # Random features
17 y = np.random.randint(2, size=num_samples) # Binary classification labels
18
19 # Create a DataFrame
20 df = pd.DataFrame(X, columns=[f'feature_{i}' for i in range(X.shape[1])])
21 df['target'] = y
22
23 # Separate features (X) and target variable (y)
24 X = df.drop(columns=['target'])
25 y = df['target']
26
27 # Standardize the data
28 scaler = StandardScaler()
29 X_scaled = scaler.fit_transform(X)
30
31 # Split the data into training and testing sets
32 X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
33
34 # Build the deep neural network model
35 model = Sequential()
36 model.add(Dense(units=64, input_dim=X_train.shape[1], activation='relu'))
37 model.add(Dense(units=32, activation='relu'))
38 model.add(Dense(units=16, activation='relu'))
39 model.add(Dense(units=1, activation='sigmoid'))
40
41 # Compile the model
42 model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
43
44 # Train the model
45 model.fit(X_train, y_train, epochs=10, batch_size=32, validation_split=0.2)
46
47 # Evaluate the model on the test set
48 y_pred_prob = model.predict(X_test)
49 y_pred = (y_pred_prob > 0.5).astype(int)
50
51 # Calculate accuracy
52 accuracy = accuracy_score(y_test, y_pred)
53 print("Test Accuracy:", accuracy)
54
```

```
Epoch 1/10
20/20 [=====] - 2s 22ms/step - loss: 0.7064 - accu
racy: 0.4922 - val_loss: 0.6907 - val_accuracy: 0.5500
Epoch 2/10
20/20 [=====] - 0s 8ms/step - loss: 0.6852 - accu
racy: 0.5359 - val_loss: 0.6877 - val_accuracy: 0.5437
Epoch 3/10
20/20 [=====] - 0s 7ms/step - loss: 0.6764 - accu
racy: 0.5922 - val_loss: 0.6870 - val_accuracy: 0.5312
Epoch 4/10
20/20 [=====] - 0s 7ms/step - loss: 0.6711 - accu
racy: 0.5953 - val_loss: 0.6866 - val_accuracy: 0.5250
Epoch 5/10
20/20 [=====] - 0s 6ms/step - loss: 0.6629 - accu
racy: 0.6109 - val_loss: 0.6858 - val_accuracy: 0.5312
Epoch 6/10
20/20 [=====] - 0s 6ms/step - loss: 0.6567 - accu
racy: 0.6328 - val_loss: 0.6860 - val_accuracy: 0.5500
Epoch 7/10
20/20 [=====] - 0s 6ms/step - loss: 0.6502 - accu
racy: 0.6375 - val_loss: 0.6842 - val_accuracy: 0.5500
Epoch 8/10
20/20 [=====] - 0s 6ms/step - loss: 0.6446 - accu
racy: 0.6406 - val_loss: 0.6889 - val_accuracy: 0.5125
Epoch 9/10
20/20 [=====] - 0s 7ms/step - loss: 0.6412 - accu
racy: 0.6438 - val_loss: 0.6918 - val_accuracy: 0.4875
Epoch 10/10
20/20 [=====] - 0s 6ms/step - loss: 0.6322 - accu
racy: 0.6750 - val_loss: 0.6894 - val_accuracy: 0.5188
7/7 [=====] - 0s 5ms/step
Test Accuracy: 0.55
```

Create a complete end to end neural network model using Keras Sequential Model and Keras Layer API

In [39]:

```
1 import numpy as np
2 import tensorflow as tf
3 from tensorflow.keras import layers, models
4 from tensorflow.keras.datasets import fashion_mnist
5 from tensorflow.keras.utils import to_categorical
6 import matplotlib.pyplot as plt
7
8 # Load and preprocess the Fashion-MNIST dataset
9 (train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
10
11 # Normalize pixel values to be between 0 and 1
12 train_images, test_images = train_images / 255.0, test_images / 255.0
13
14 # One-hot encode the labels
15 train_labels = to_categorical(train_labels)
16 test_labels = to_categorical(test_labels)
17
18 # Build the neural network using Keras Sequential Model and Keras Layer
19 model = models.Sequential()
20
21 # Flatten layer to flatten the input
22 model.add(layers.Flatten(input_shape=(28, 28)))
23
24 # Dense layers with ReLU activation
25 model.add(layers.Dense(128, activation='relu'))
26 model.add(layers.Dense(64, activation='relu'))
27
28 # Output layer with softmax activation for multi-class classification
29 model.add(layers.Dense(10, activation='softmax'))
30
31 # Compile the model
32 model.compile(optimizer='adam',
33               loss='categorical_crossentropy',
34               metrics=['accuracy'])
35
36 # Display the summary of the model
37 model.summary()
38
39 # Train the model
40 history = model.fit(train_images, train_labels, epochs=10, validation_data=(test_images, test_labels))
41
42 # Evaluate the model on the test set
43 test_loss, test_accuracy = model.evaluate(test_images, test_labels)
44 print(f'Test Accuracy: {test_accuracy}')
45
46 # Plot training history (optional)
47 plt.figure(figsize=(12, 4))
48
49 plt.subplot(1, 2, 1)
50 plt.plot(history.history['accuracy'], label='Training Accuracy')
51 plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
52 plt.xlabel('Epochs')
53 plt.ylabel('Accuracy')
54 plt.legend()
55
56 plt.subplot(1, 2, 2)
57 plt.plot(history.history['loss'], label='Training Loss')
58 plt.plot(history.history['val_loss'], label='Validation Loss')
59 plt.xlabel('Epochs')
60 plt.ylabel('Loss')
61 plt.legend()
```

```
62  
63 plt.tight_layout()  
64 plt.show()  
65
```

Epoch 9/10


1500/1500 [=====] - 9s 6ms/step - loss: 0.2475
- accuracy: 0.9071 - val_loss: 0.3222 - val_accuracy: 0.8853

Epoch 10/10

1500/1500 [=====] - 8s 5ms/step - loss: 0.2384
- accuracy: 0.9115 - val_loss: 0.3303 - val_accuracy: 0.8869

313/313 [=====] - 1s 3ms/step - loss: 0.3461 -
accuracy: 0.8807

Test Accuracy: 0.8806999921798706



MNIST dataset (classify handwritten numerals)

or fashion-MNIST dataset or dataset from other source

This code defines a simple convolutional neural network (CNN) using the Keras API with TensorFlow backend. It loads the MNIST dataset, preprocesses the data, builds the model, compiles it, trains it, and evaluates its performance on the test set.

```

In [21]: 1 import tensorflow as tf
2 from tensorflow.keras import layers, models
3 from tensorflow.keras.datasets import mnist
4 from tensorflow.keras.utils import to_categorical
5
6 # Load and preprocess the MNIST dataset
7 (train_images, train_labels), (test_images, test_labels) = mnist.load_data
8 train_images = train_images.reshape((60000, 28, 28, 1)).astype('float32')
9 test_images = test_images.reshape((10000, 28, 28, 1)).astype('float32')
10 train_labels = to_categorical(train_labels)
11 test_labels = to_categorical(test_labels)
12
13 # Build a simple convolutional neural network (CNN)
14 model = models.Sequential()
15 model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28,
16 model.add(layers.MaxPooling2D((2, 2)))
17 model.add(layers.Conv2D(64, (3, 3), activation='relu'))
18 model.add(layers.MaxPooling2D((2, 2)))
19 model.add(layers.Conv2D(64, (3, 3), activation='relu'))
20 model.add(layers.Flatten())
21 model.add(layers.Dense(64, activation='relu'))
22 model.add(layers.Dense(10, activation='softmax'))
23
24 # Compile the model
25 model.compile(optimizer='adam',
26               loss='categorical_crossentropy',
27               metrics=['accuracy'])
28
29 # Train the model
30 model.fit(train_images, train_labels, epochs=5, batch_size=64, validation_data=(test_images, test_labels))
31
32 # Evaluate the model on the test set
33 test_loss, test_acc = model.evaluate(test_images, test_labels)
34 print(f'Test accuracy: {test_acc}')
35
36

```

Epoch 1/5

750/750 [=====] - 60s 77ms/step - loss: 0.2014 - accuracy: 0.9379 - val_loss: 0.0767 - val_accuracy: 0.9764

Epoch 2/5

750/750 [=====] - 59s 78ms/step - loss: 0.0564 - accuracy: 0.9827 - val_loss: 0.0551 - val_accuracy: 0.9818

Epoch 3/5

750/750 [=====] - 50s 67ms/step - loss: 0.0391 - accuracy: 0.9873 - val_loss: 0.0470 - val_accuracy: 0.9863

Epoch 4/5

750/750 [=====] - 55s 73ms/step - loss: 0.0297 - accuracy: 0.9903 - val_loss: 0.0448 - val_accuracy: 0.9872

Epoch 5/5

750/750 [=====] - 61s 82ms/step - loss: 0.0235 - accuracy: 0.9920 - val_loss: 0.0416 - val_accuracy: 0.9895

313/313 [=====] - 4s 12ms/step - loss: 0.0318 - accuracy: 0.9893

Test accuracy: 0.989300012588501

```
In [18]: 1 from sklearn import datasets
        2
```

```
In [2]: 1 digits = datasets.load_digits()
        2 dir(digits)
```

```
Out[2]: ['DESCR', 'data', 'feature_names', 'frame', 'images', 'target', 'target_names']
```

```
In [3]: 1 print(digits.images[0])

[[ 0.  0.  5. 13.  9.  1.  0.  0.]
 [ 0.  0. 13. 15. 10. 15.  5.  0.]
 [ 0.  3. 15.  2.  0. 11.  8.  0.]
 [ 0.  4. 12.  0.  0.  8.  8.  0.]
 [ 0.  5.  8.  0.  0.  9.  8.  0.]
 [ 0.  4. 11.  0.  1. 12.  7.  0.]
 [ 0.  2. 14.  5. 10. 12.  0.  0.]
 [ 0.  0.  6. 13. 10.  0.  0.  0.]]
```

```
In [5]: 1 import matplotlib.pyplot as plt
        2 def plot_multi(i):
        3     nplots = 16
        4     fig = plt.figure(figsize=(15, 15))
        5     for j in range(nplots):
        6         plt.subplot(4, 4, j+1)
        7         plt.imshow(digits.images[i+j], cmap='binary')
        8         plt.title(digits.target[i+j])
        9         plt.axis('off')
       10         plt.show()
       11         plot_multi(0)
       12 y = digits.target
       13 x = digits.images.reshape((len(digits.images), -1))
       14 # gives the shape of the data
       15 x.shape
       16
```

```
Out[5]: (1797, 64)
```

```
In [6]: 1 x[0]
```

```
Out[6]: array([ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.,  0.,  0., 13., 15., 10.,
                15.,  5.,  0.,  0.,  3., 15.,  2.,  0., 11.,  8.,  0.,  0.,  4.,
                12.,  0.,  0.,  8.,  8.,  0.,  0.,  5.,  8.,  0.,  0.,  9.,  8.,
                0.,  0.,  4., 11.,  0.,  1., 12.,  7.,  0.,  0.,  2., 14.,  5.,
                10., 12.,  0.,  0.,  0.,  0.,  6., 13., 10.,  0.,  0.,  0.]
```

```
In [7]: 1 x_train = x[:1000]
        2 y_train = y[:1000]
        3 x_test = x[1000:]
        4 y_test = y[1000:]
        5
```

In [8]: 1 x_train

```
Out[8]: array([[ 0.,  0.,  5., ...,  0.,  0.,  0.],
               [ 0.,  0.,  0., ..., 10.,  0.,  0.],
               [ 0.,  0.,  0., ..., 16.,  9.,  0.],
               ...,
               [ 0.,  0.,  3., ...,  9.,  0.,  0.],
               [ 0.,  0.,  0., ...,  6.,  0.,  0.],
               [ 0.,  0.,  9., ..., 10.,  0.,  0.]])
```


In [11]: 1 x_test

```
Out[11]: array([[ 0.,  0.,  1., ..., 15., 16., 15.],
                 [ 0.,  0.,  0., ...,  5.,  0.,  0.],
                 [ 0.,  0.,  6., ...,  3.,  0.,  0.],
                 ...,
                 [ 0.,  0.,  1., ...,  6.,  0.,  0.],
                 [ 0.,  0.,  2., ..., 12.,  0.,  0.],
                 [ 0.,  0., 10., ..., 12.,  1.,  0.]])
```

```
In [12]: 1 from sklearn.neural_network import MLPClassifier
2 mlp = MLPClassifier(hidden_layer_sizes=(15,),
3   activation='logistic',
4   alpha=1e-4, solver='sgd',
5   tol=1e-4, random_state=1,
6   learning_rate_init=.1,
7   verbose=True)
8 mlp.fit(x_train, y_train)
9
```

```
Iteration 1, loss = 2.22958289
Iteration 2, loss = 1.91207743
Iteration 3, loss = 1.62507727
Iteration 4, loss = 1.32649842
Iteration 5, loss = 1.06100535
Iteration 6, loss = 0.83995513
Iteration 7, loss = 0.67806075
Iteration 8, loss = 0.55175832
Iteration 9, loss = 0.45840445
Iteration 10, loss = 0.39149735
Iteration 11, loss = 0.33676351
Iteration 12, loss = 0.29059880
Iteration 13, loss = 0.25437208
Iteration 14, loss = 0.22838372
Iteration 15, loss = 0.20200554
Iteration 16, loss = 0.18186565
Iteration 17, loss = 0.16461183
Iteration 18, loss = 0.14990228
Iteration 19, loss = 0.13892154
Iteration 20, loss = 0.12833784
```

```
In [13]: 1 fig, axes = plt.subplots(1,1)
2 axes.plot(mlp.loss_curve_, 'o-')
3 axes.set_xlabel("number of iteration")
4 axes.set_ylabel("loss")
5 plt.show()
6
```



```
In [15]: 1 predictions = mlp.predict(x_test)
2 predictions[:50]
3
```

```
Out[15]: array([1, 4, 0, 5, 3, 6, 9, 6, 1, 7, 5, 4, 4, 7, 2, 8, 2, 2, 5, 7, 9, 5,
4, 4, 9, 0, 8, 9, 8, 0, 1, 2, 3, 4, 5, 6, 7, 8, 3, 0, 1, 2, 3, 4,
5, 6, 7, 8, 5, 0])
```

```
In [16]: 1 y_test[:50]
2
```

```
Out[16]: array([1, 4, 0, 5, 3, 6, 9, 6, 1, 7, 5, 4, 4, 7, 2, 8, 2, 2, 5, 7, 9, 5,
4, 4, 9, 0, 8, 9, 8, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 0, 1, 2, 3, 4,
5, 6, 7, 8, 9, 0])
```

```
In [17]: 1 from sklearn.metrics import accuracy_score
2 accuracy_score(y_test, predictions)
```

```
Out[17]: 0.9146800501882058
```

```
In [ ]: 1
```

