

```
In [1]: 1 import pandas as pd
2 from sklearn.model_selection import train_test_split
3 from keras.models import Sequential
4 from keras.layers import Activation,Dense
```

C:\Users\Anusha V\anaconda3\lib\site-packages\scipy__init__.py:155: UserWarning: A NumPy version >=1.18.5 and <1.25.0 is required for this version of SciPy (detected version 1.26.1
 warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")

```
In [28]: 1
2 data=pd.read_csv(r"C:\Users\Anusha V\Downloads\heart1.csv")
```

```
In [29]: 1 data
2
```

```
Out[29]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	t
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	
...	
1020	59	1	1	140	221	0	1	164	1	0.0	2	0	2	
1021	60	1	0	125	258	0	0	141	1	2.8	1	1	3	
1022	47	1	0	110	275	0	0	118	1	1.0	1	1	2	
1023	50	0	0	110	254	0	0	159	0	0.0	2	0	2	
1024	54	1	0	120	188	0	1	113	0	1.4	1	1	3	

1025 rows × 14 columns



```
In [30]: 1 X = data.drop(columns=['target'])
2 y = data['target']
3
```

```
In [31]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2
2
```

```
In [32]: 1 model = Sequential()
2 model.add(Dense(32, activation='relu', input_shape=(X_train.shape[1],)))
3 model.add(Dense(16, activation='relu'))
4 model.add(Dense(1, activation='sigmoid'))
5
```

```
In [ ]: 1
2 model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['a
3
4
5 model.fit(X_train, y_train, epochs=20, batch_size=32, validation_split=
6
7
```

Epoch 1/20

21/21 [=====] - 2s 21ms/step - loss: 4.2801 - accuracy: 0.4649 - val_loss: 1.1501 - val_accuracy: 0.4390

Epoch 2/20

21/21 [=====] - 0s 11ms/step - loss: 1.1572 - accuracy: 0.5091 - val_loss: 0.8599 - val_accuracy: 0.5183

Epoch 3/20

21/21 [=====] - 0s 9ms/step - loss: 0.8155 - accuracy: 0.6098 - val_loss: 0.8361 - val_accuracy: 0.5488

Epoch 4/20

21/21 [=====] - 0s 9ms/step - loss: 0.7127 - accuracy: 0.6463 - val_loss: 0.8186 - val_accuracy: 0.6037

Epoch 5/20

21/21 [=====] - 0s 9ms/step - loss: 0.6476 - accuracy: 0.6616 - val_loss: 0.8170 - val_accuracy: 0.5976

Epoch 6/20

21/21 [=====] - 0s 9ms/step - loss: 0.6019 - accuracy: 0.6905 - val_loss: 0.7779 - val_accuracy: 0.6280

Epoch 7/20

21/21 [=====] - 0s 8ms/step - loss: 0.5743 - accuracy: 0.7180 - val_loss: 0.7930 - val_accuracy: 0.6037

Epoch 8/20

21/21 [=====] - 0s 10ms/step - loss: 0.5606 - accuracy: 0.7287 - val_loss: 0.7741 - val_accuracy: 0.6159

Epoch 9/20

21/21 [=====] - 0s 8ms/step - loss: 0.5435 - accuracy: 0.7363 - val_loss: 0.7611 - val_accuracy: 0.6220

Epoch 10/20

21/21 [=====] - 0s 11ms/step - loss: 0.5296 - accuracy: 0.7409 - val_loss: 0.7394 - val_accuracy: 0.6280

Epoch 11/20

21/21 [=====] - 0s 8ms/step - loss: 0.5383 - accuracy: 0.7317 - val_loss: 0.7465 - val_accuracy: 0.6463

Epoch 12/20

21/21 [=====] - 0s 8ms/step - loss: 0.5158 - accuracy: 0.7332 - val_loss: 0.6861 - val_accuracy: 0.6402

Epoch 13/20

21/21 [=====] - 0s 10ms/step - loss: 0.5326 - accuracy: 0.7363 - val_loss: 0.7073 - val_accuracy: 0.6524

Epoch 14/20

21/21 [=====] - 0s 9ms/step - loss: 0.4983 - accuracy: 0.7622 - val_loss: 0.6782 - val_accuracy: 0.6280

Epoch 15/20

21/21 [=====] - 0s 9ms/step - loss: 0.4793 - accuracy: 0.7744 - val_loss: 0.6642 - val_accuracy: 0.6402

Epoch 16/20

21/21 [=====] - 0s 8ms/step - loss: 0.4748 - accuracy: 0.7759 - val_loss: 0.6361 - val_accuracy: 0.6524

Epoch 17/20

13/21 [=====>.....] - ETA: 0s - loss: 0.4944 - accuracy: 0.7500

```
In [ ]: 1 test_loss, test_acc = model.evaluate(X_test, y_test)
        2 print('Test accuracy:', test_acc)
```

```
In [ ]: 1
```

```
In [ ]: 1 import pandas as pd
        2 from sklearn.model_selection import train_test_split
        3 from sklearn.preprocessing import StandardScaler, OneHotEncoder
        4 from sklearn.impute import SimpleImputer
        5 from sklearn.compose import ColumnTransformer
        6 from sklearn.pipeline import Pipeline
        7 from imblearn.over_sampling import SMOTE
        8
```

```
In [ ]: 1 # Assuming your dataset is in a CSV file
        2 df = pd.read_csv(r"C:\Users\Anusha V\Desktop\diabetes.csv")
        3
```

```
In [ ]: 1 # Check for missing values
        2 print(df.isnull().sum())
        3
        4 # Check data types
        5 print(df.dtypes)
        6
```

```
In [ ]: 1 data.head(2)
```

```
In [ ]: 1 X = data.drop(columns=['count'])
        2 y = data['count']
```

```
In [24]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2
        2
```

```
In [25]: 1 model = Sequential()
        2 model.add(Dense(32, activation='relu', input_shape=(X_train.shape[1],)))
        3 model.add(Dense(16, activation='relu'))
        4 model.add(Dense(1, activation='sigmoid'))
        5
```

```
In [26]: 1 model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['a
2
3
4 model.fit(X_train, y_train, epochs=20, batch_size=32, validation_split=
5
6
```

```
Epoch 1/20
5/5 [=====] - 3s 182ms/step - loss: -0.4637 - accuracy: 0.1375 - val_loss: -5.4125 - val_accuracy: 0.1000
Epoch 2/20
5/5 [=====] - 0s 63ms/step - loss: -1.1090 - accuracy: 0.1375 - val_loss: -7.5187 - val_accuracy: 0.1000
Epoch 3/20
5/5 [=====] - 0s 33ms/step - loss: -1.6182 - accuracy: 0.1375 - val_loss: -9.8189 - val_accuracy: 0.1000
Epoch 4/20
5/5 [=====] - 0s 30ms/step - loss: -2.2691 - accuracy: 0.1375 - val_loss: -12.1385 - val_accuracy: 0.1000
Epoch 5/20
5/5 [=====] - 0s 34ms/step - loss: -2.9196 - accuracy: 0.1375 - val_loss: -14.5659 - val_accuracy: 0.1000
Epoch 6/20
5/5 [=====] - 0s 34ms/step - loss: -3.5364 - accuracy: 0.1375 - val_loss: -17.2265 - val_accuracy: 0.1000
Epoch 7/20
5/5 [=====] - 0s 27ms/step - loss: -4.2658 - accuracy: 0.1375 - val_loss: -19.9719 - val_accuracy: 0.1000
Epoch 8/20
5/5 [=====] - 0s 24ms/step - loss: -5.0590 - accuracy: 0.1375 - val_loss: -22.8575 - val_accuracy: 0.1000
Epoch 9/20
5/5 [=====] - 0s 19ms/step - loss: -5.8365 - accuracy: 0.1375 - val_loss: -25.9913 - val_accuracy: 0.1000
Epoch 10/20
5/5 [=====] - 0s 23ms/step - loss: -6.7541 - accuracy: 0.1375 - val_loss: -29.3319 - val_accuracy: 0.1000
Epoch 11/20
5/5 [=====] - 0s 19ms/step - loss: -7.7242 - accuracy: 0.1375 - val_loss: -33.0192 - val_accuracy: 0.1000
Epoch 12/20
5/5 [=====] - 0s 37ms/step - loss: -8.6385 - accuracy: 0.1375 - val_loss: -37.1766 - val_accuracy: 0.1000
Epoch 13/20
5/5 [=====] - 0s 33ms/step - loss: -9.8316 - accuracy: 0.1437 - val_loss: -41.3374 - val_accuracy: 0.1250
Epoch 14/20
5/5 [=====] - 0s 20ms/step - loss: -10.9624 - accuracy: 0.1500 - val_loss: -45.8264 - val_accuracy: 0.1250
Epoch 15/20
5/5 [=====] - 0s 20ms/step - loss: -12.1072 - accuracy: 0.1500 - val_loss: -50.7634 - val_accuracy: 0.1250
Epoch 16/20
5/5 [=====] - 0s 27ms/step - loss: -13.5395 - accuracy: 0.1625 - val_loss: -55.8706 - val_accuracy: 0.1250
Epoch 17/20
5/5 [=====] - 0s 24ms/step - loss: -14.8398 - accuracy: 0.1688 - val_loss: -61.6153 - val_accuracy: 0.1250
Epoch 18/20
5/5 [=====] - 0s 29ms/step - loss: -16.1310 - accuracy: 0.1688 - val_loss: -68.0390 - val_accuracy: 0.1250
Epoch 19/20
5/5 [=====] - 0s 23ms/step - loss: -17.9924 - accuracy: 0.1688 - val_loss: -74.4900 - val_accuracy: 0.1250
Epoch 20/20
5/5 [=====] - 0s 25ms/step - loss: -19.7415 - accuracy: 0.1688 - val_loss: -81.4427 - val_accuracy: 0.1250
```

Out[26]: <keras.src.callbacks.History at 0x1fd40accfa0>

```
In [27]: 1 test_loss, test_acc = model.evaluate(X_test, y_test)
          2 print('Test accuracy:', test_acc)
```

2/2 [=====] - 0s 18ms/step - loss: -29.0556 - accuracy: 0.1600
Test accuracy: 0.1599999964237213

Through a step-by-step process calculate TF/IDF for the given corpus and mention the words having highest value

Doc1: we are going to Mysore Doc2: Mysore is a famous place Doc3: we are going to famous place

```

In [2]: 1 from sklearn.feature_extraction.text import TfidfVectorizer
2
3 # Define the corpus
4 corpus = [
5     "we are going to Mysore",
6     "Mysore is a famous place",
7     "we are going to famous place"
8 ]
9
10 #Create a TF-IDF Vectorizer
11 tfidf_vectorizer = TfidfVectorizer()
12
13 # Fit and transform the corpus
14 tfidf_matrix = tfidf_vectorizer.fit_transform(corpus)
15
16 # Get feature names (words)
17 feature_names = tfidf_vectorizer.get_feature_names_out()
18
19 #Create a dictionary to store the TF-IDF values for each word in each d
20 tfidf_values = {}
21
22 #Loop through each document and each word to get the TF-IDF value
23 for doc_index, doc in enumerate(corpus):
24     feature_index = tfidf_matrix[doc_index, :].nonzero()[1]
25     tfidf_doc = zip(feature_index, [tfidf_matrix[doc_index, x] for x in
26
27         for word_index, tfidf in tfidf_doc:
28             word = feature_names[word_index]
29             if word not in tfidf_values:
30                 tfidf_values[word] = [(doc_index, tfidf)]
31             else:
32                 tfidf_values[word].append((doc_index, tfidf))
33
34 #Find the words with the highest TF-IDF values
35 highest_tfidf_words = {}
36 for word, values in tfidf_values.items():
37     highest_tfidf_words[word] = max(values, key=lambda x: x[1])
38
39 # Print the words with the highest TF-IDF values
40 for word, (doc_index, tfidf) in highest_tfidf_words.items():
41     print(f"Word: {word}, Document: {doc_index + 1}, TF-IDF Value: {tfi
42

```

C:\Users\Anusha V\anaconda3\lib\site-packages\scipy__init__.py:155: UserWarning: A NumPy version >=1.18.5 and <1.25.0 is required for this version of SciPy (detected version 1.26.2)

warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")

Word: mysore, Document: 2, TF-IDF Value: 0.4598535287588349
 Word: to, Document: 1, TF-IDF Value: 0.4472135954999579
 Word: going, Document: 1, TF-IDF Value: 0.4472135954999579
 Word: are, Document: 1, TF-IDF Value: 0.4472135954999579
 Word: we, Document: 1, TF-IDF Value: 0.4472135954999579
 Word: place, Document: 2, TF-IDF Value: 0.4598535287588349
 Word: famous, Document: 2, TF-IDF Value: 0.4598535287588349
 Word: is, Document: 2, TF-IDF Value: 0.6046521283053111

Create a complete end to end Neural network for MNIST(classify handwritten numerals)

In [3]:

```
1  # Import necessary Libraries
2  import tensorflow as tf
3  from tensorflow.keras import layers, models
4  from tensorflow.keras.datasets import mnist
5  from tensorflow.keras.utils import to_categorical
6
7  # Load and preprocess the MNIST dataset
8  (train_images, train_labels), (test_images, test_labels) = mnist.load_d
9
10 # Normalize pixel values to be between 0 and 1
11 train_images, test_images = train_images / 255.0, test_images / 255.0
12
13 # One-hot encode the Labels
14 train_labels = to_categorical(train_labels)
15 test_labels = to_categorical(test_labels)
16
17 # Build the neural network model
18 model = models.Sequential()
19 model.add(layers.Flatten(input_shape=(28, 28))) # Flatten the 28x28 im
20 model.add(layers.Dense(128, activation='relu')) # Hidden Layer with 12
21 model.add(layers.Dropout(0.2)) # Dropout Layer to reduce overfitting
22 model.add(layers.Dense(10, activation='softmax')) # Output Layer with
23
24 # Compile the model
25 model.compile(optimizer='adam',
26               loss='categorical_crossentropy',
27               metrics=['accuracy'])
28
29 # Train the model
30 model.fit(train_images, train_labels, epochs=5, batch_size=64, validati
31
32 # Evaluate the model on the test set
33 test_loss, test_acc = model.evaluate(test_images, test_labels)
34 print(f"Test Accuracy: {test_acc}")
35
36 # Make predictions on a few test images
37 predictions = model.predict(test_images[:5])
38 predicted_labels = tf.argmax(predictions, axis=1)
39
40 # Print the predicted Labels
41 print("Predicted Labels:", predicted_labels.numpy())
42
```

```

Epoch 1/5
938/938 [=====] - 16s 14ms/step - loss: 0.3399 - accuracy: 0.9026 - val_loss: 0.1606 - val_accuracy: 0.9515
Epoch 2/5
938/938 [=====] - 10s 10ms/step - loss: 0.1634 - accuracy: 0.9514 - val_loss: 0.1130 - val_accuracy: 0.9658
Epoch 3/5
938/938 [=====] - 7s 7ms/step - loss: 0.1201 - accuracy: 0.9644 - val_loss: 0.0899 - val_accuracy: 0.9725
Epoch 4/5
938/938 [=====] - 8s 8ms/step - loss: 0.0981 - accuracy: 0.9704 - val_loss: 0.0831 - val_accuracy: 0.9747
Epoch 5/5
938/938 [=====] - 7s 7ms/step - loss: 0.0827 - accuracy: 0.9750 - val_loss: 0.0765 - val_accuracy: 0.9772
313/313 [=====] - 1s 4ms/step - loss: 0.0765 - accuracy: 0.9772
Test Accuracy: 0.9771999716758728
1/1 [=====] - 0s 288ms/step
Predicted Labels: [7 2 1 0 4]

```

N-grams are defined as the combinations of N Keywords together consider the given

“The greatest glory in Living lies not in never falling but in raising every Lies” Generate bi grams for the above Generate Tri grams for the above

```

In [4]: 1 from nltk import ngrams
        2 from nltk.tokenize import word_tokenize
        3
        4 # Given text
        5 text = "The greatest glory in Living lies not in never falling but in r
        6
        7 # Tokenize the text into words
        8 words = word_tokenize(text)
        9
       10 # Function to generate n-grams
       11 def generate_ngrams(tokens, n):
       12     n_grams = ngrams(tokens, n)
       13     return [' '.join(gram) for gram in n_grams]
       14
       15 # Generate bi-grams
       16 bi_grams = generate_ngrams(words, 2)
       17 print("Bi-grams:", bi_grams)
       18
       19 # Generate tri-grams
       20 tri_grams = generate_ngrams(words, 3)
       21 print("Tri-grams:", tri_grams)
       22

```

```

Bi-grams: ['The greatest', 'greatest glory', 'glory in', 'in Living', 'Living lies', 'lies not', 'not in', 'in never', 'never falling', 'falling but', 'but in', 'in raising', 'raising every', 'every Lies']
Tri-grams: ['The greatest glory', 'greatest glory in', 'glory in Living', 'in Living lies', 'Living lies not', 'lies not in', 'not in never', 'in never falling', 'never falling but', 'falling but in', 'but in raising', 'in raising every', 'raising every Lies']

```

In []:

1