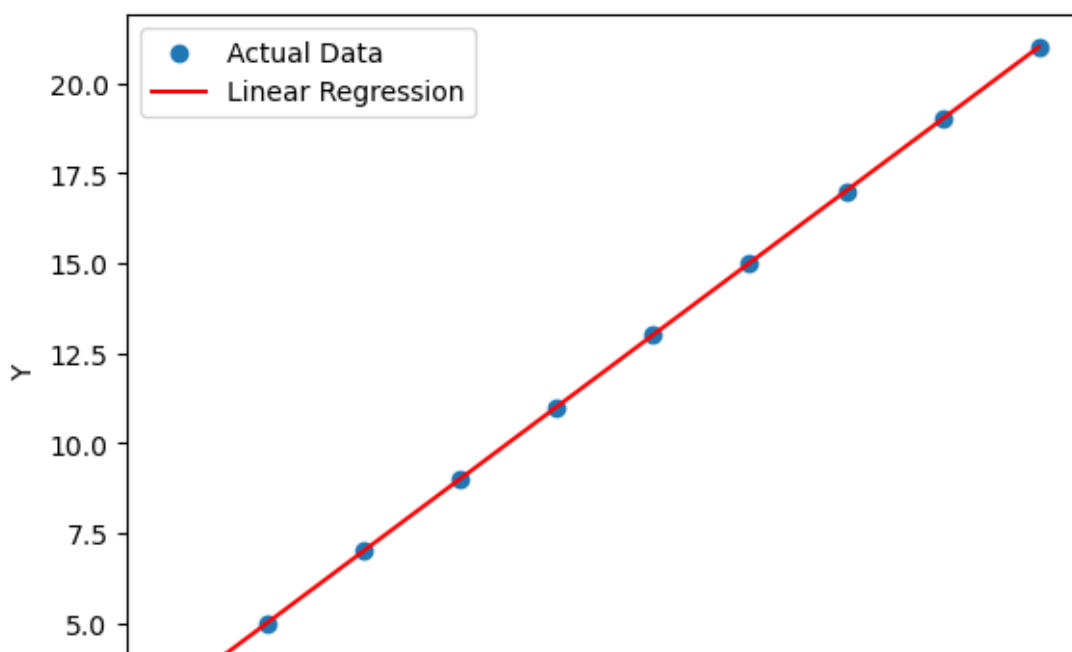


Simple Linear Regression

In [14]:

```
1 import numpy as np
2 from sklearn.linear_model import LinearRegression
3 import matplotlib.pyplot as plt
4
5 # Generate synthetic data
6 X = np.arange(1, 11).reshape(-1, 1)
7 y = 2 * X + 1
8
9 # Create and train a simple linear regression model
10 model = LinearRegression()
11 model.fit(X, y)
12
13 # Make predictions
14 y_pred = model.predict(X)
15
16 # Plot the data and regression line
17 plt.scatter(X, y, label="Actual Data")
18 plt.plot(X, y_pred, color='red', label="Linear Regression")
19 plt.xlabel("X")
20 plt.ylabel("Y")
21 plt.legend()
22 plt.show()
```




Multiple Linear Regression

```
In [15]: 1 # Generate synthetic data with two features
2 X = np.array([[1, 2], [2, 3], [3, 4], [4, 5], [5, 6]])
3 y = 2 * X[:, 0] + 3 * X[:, 1] + 1
4
5 # Create and train a multiple linear regression model
6 model = LinearRegression()
7 model.fit(X, y)
8
9 # Make predictions
10 y_pred = model.predict(X)
11 y_pred
```

```
Out[15]: array([ 9., 14., 19., 24., 29.])
```

Polynomial Regression:

```
In [16]: 1 from sklearn.preprocessing import PolynomialFeatures
2 from sklearn.pipeline import make_pipeline
3
4 # Generate synthetic data
5 X = np.arange(1, 11).reshape(-1, 1)
6 y = 0.5 * X**2 - 2 * X + 1
7
8 # Create a polynomial regression model
9 degree = 2 # Degree of the polynomial
10 polyreg = make_pipeline(PolynomialFeatures(degree), LinearRegression())
11 polyreg.fit(X, y)
12
13 # Make predictions
14 y_pred = polyreg.predict(X)
15
16 # Plot the data and polynomial regression curve
17 plt.scatter(X, y, label="Actual Data")
18 plt.plot(X, y_pred, color='red', label=f"Polynomial Regression (Degree
19 plt.xlabel("X")
20 plt.ylabel("Y")
21 plt.legend()
22 plt.show()
```



Evaluation Matrix

R Squared[R2]

```
In [17]: 1 from sklearn.metrics import mean_absolute_error
2 # Actual values (ground truth)
3 actual = [10, 20, 30, 40, 50]
4
5 # Predicted values from your model
6 predicted = [12, 18, 28, 37, 45]
7
8 # Calculate the Mean Absolute Error
9 mae = mean_absolute_error(actual, predicted)
10
11 print(f"Mean Absolute Error: {mae}")
```

Mean Absolute Error: 2.8

mean_absolute_error

```
In [1]: 1 import numpy as np
2 from sklearn.metrics import mean_squared_error
3 from math import sqrt
4
5 # Actual values
6 actual = [10, 12, 15, 18, 20]
7
8 # Predicted values
9 predicted = [9, 11, 14, 17, 19]
10
11 # Calculate Mean Squared Error (MSE)
12 mse = mean_squared_error(actual, predicted)
13
14 # Calculate RMSE
15 rmse = sqrt(mse)
16
17 print("RMSE:", rmse)
```

RMSE: 1.0

R-squared (R2) score

```
In [19]: 1 from sklearn.metrics import r2_score
2
3 # Assuming you have observed actual and predicted values
4 actual_values = [10, 20, 30, 40, 50]
5 predicted_values = [12, 18, 28, 38, 52]
6
7 # Calculate the R-squared (R2) score
8 r2 = r2_score(actual_values, predicted_values)
9
10 print(f"R-squared (R2) Score: {r2}")
```

R-squared (R2) Score: 0.98

Cross Validation

Hold-Out Validation:

```
In [20]: 1 from sklearn.model_selection import train_test_split
2 from sklearn.datasets import load_iris # You can replace this with your own data
3
4 # Load your dataset (replace with your data loading code)
5 data = load_iris()
6
7 # Split the dataset into features (X) and target labels (y)
8 X = data.data
9 y = data.target
10
11 # Split the data into a training set and a test set
12 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
13 X
```

```
In [21]: 1 from sklearn.model_selection import LeaveOneOut
2 from sklearn.linear_model import LinearRegression
3 loo = LeaveOneOut()
4 total_mse = 0
5
6 for train_index, test_index in loo.split(X):
7     X_train, X_test = X[train_index], X[test_index]
8     y_train, y_test = y[train_index], y[test_index]
9
10    model = LinearRegression()
11    model.fit(X_train, y_train)
12    y_pred = model.predict(X_test)
13
14    mse = mean_squared_error(y_test, y_pred)
15    total_mse += mse
16
17 average_mse = total_mse / len(X)
18 print("Average Mean Squared Error (LOOCV):", average_mse)
```

Average Mean Squared Error (LOOCV): 0.04954185325184582

K-Fold Cross-Validation:

```
In [22]: 1 from sklearn.model_selection import KFold
2
3 k = 5 # You can choose the number of folds
4 kf = KFold(n_splits=k)
5 total_mse = 0
6
7 for train_index, test_index in kf.split(X):
8     X_train, X_test = X[train_index], X[test_index]
9     y_train, y_test = y[train_index], y[test_index]
10
11    model = LinearRegression()
12    model.fit(X_train, y_train)
13    y_pred = model.predict(X_test)
14
15    mse = mean_squared_error(y_test, y_pred)
16    total_mse += mse
17
18 average_mse = total_mse / k
19 print(f"Average Mean Squared Error ({k}-fold CV):", average_mse)
```

Average Mean Squared Error (5-fold CV): 0.06897012554624687

```
In [4]: 1 from sklearn.model_selection import KFold
2 from sklearn.model_selection import train_test_split
3 from sklearn.metrics import accuracy_score
4 import numpy as np
5
6 # Generate some example data
7 X = np.array([[1, 2], [2, 3], [3, 4], [4, 5], [5, 6], [6, 7]])
8 y = np.array([0, 0, 1, 1, 2, 2])
9
10 # Specify the number of folds
11 k = 3
12
13 # Initialize the KFold class
14 kf = KFold(n_splits=k, shuffle=True, random_state=42)
15
16 # Initialize a variable to keep track of the average accuracy
17 avg_accuracy = 0
18
19 # Loop through each fold
20 for train_index, test_index in kf.split(X):
21     X_train, X_test = X[train_index], X[test_index]
22     y_train, y_test = y[train_index], y[test_index]
23
24     # Train your machine Learning model on the training data
25     # For simplicity, let's use a basic classifier (e.g., Decision Tree)
26     from sklearn.tree import DecisionTreeClassifier
27     clf = DecisionTreeClassifier()
28     clf.fit(X_train, y_train)
29
30     # Make predictions on the test data
31     y_pred = clf.predict(X_test)
32
33     # Calculate the accuracy for this fold
34     accuracy = accuracy_score(y_test, y_pred)
35     avg_accuracy += accuracy
36
37 # Calculate the average accuracy across all folds
38 avg_accuracy /= k
39
40 print("Average Accuracy:", avg_accuracy)
41
```

Average Accuracy: 0.5

```
In [8]: 1 from sklearn.model_selection import KFold
2 from sklearn.tree import DecisionTreeClassifier
3 from sklearn.metrics import accuracy_score
4 import numpy as np
5
6 X = np.array([[1, 2], [2, 3], [3, 4], [4, 5], [5, 6], [6, 7]])
7 y = np.array([0, 0, 1, 1, 2, 2])
8 k = 3
9 kf = KFold(n_splits=k, shuffle=True, random_state=42)
10
11 avg_accuracy = sum(accuracy_score(y[test], DecisionTreeClassifier().fit
12
13 print("Average Accuracy:", avg_accuracy)
14
```

Average Accuracy: 0.5

Regression Algorithm

Random Forest Regressor

```
In [29]: 1 from sklearn.ensemble import RandomForestRegressor
2 from sklearn.datasets import make_regression
3 X, y = make_regression(n_features=4, n_informative=2, random_state=0, s
4 rfr = RandomForestRegressor(max_depth=3)
5 rfr.fit(X, y)
6 print(rfr.predict([[0, 1, 0, 1]]))
```

[33.07556472]

Lasso Regression

```
In [30]: 1 from sklearn import linear_model
2 import numpy as np
3 rng = np.random.RandomState(1)
4 X = np.sort(5 * rng.rand(80, 1), axis=0)
5 y = np.sin(X).ravel()
6 y[::5] += 3 * (0.5 - rng.rand(16))
7 # Fit regression model
8 lassoReg = linear_model.Lasso(alpha=0.1)
9 lassoReg.fit(X,y)
10 # Predict
11 X_test = np.arange(0.0, 5.0, 1)[: , np.newaxis]
12 lassoReg.predict(X_test)
```

Out[30]: array([0.78305084, 0.49957596, 0.21610108, -0.0673738 , -0.35084868])

Support Vector Regression


```
In [31]: 1 from sklearn.svm import SVR
2 import numpy as np
3 rng = np.random.RandomState(1)
4 X = np.sort(5 * rng.rand(80, 1), axis=0)
5 y = np.sin(X).ravel()
6 y[::5] += 3 * (0.5 - rng.rand(16))
7 # Fit regression model
8 svr = SVR().fit(X, y)
9 # Predict
10 X_test = np.arange(0.0, 5.0, 1)[: , np.newaxis]
11 svr.predict(X_test)
```

Out[31]: array([-0.07840308, 0.78077042, 0.81326895, 0.08638149, -0.6928019])

TypesOfClassifications

Binary Classification

```
In [40]: 1 from sklearn.model_selection import train_test_split
2 from sklearn.linear_model import LogisticRegression
3
4 X = [[1.2, 2.3], [2.4, 3.5], [3.6, 4.7], [4.8, 5.9], [5.0, 6.1]]
5 y = [0, 1, 0, 1, 1]
6 # Load and preprocess your data
7 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
8
9 # Create and train a binary classification model (e.g., logistic regres
10 model = LogisticRegression()
11 model.fit(X_train, y_train)
12
13
14 # Make predictions
15 predictions = model.predict(X_test)
16 predictions
```

Out[40]: array([1])

multi-class classification

```
In [62]: 1 from sklearn.model_selection import train_test_split
2 from sklearn.svm import SVC
3
4 # Load and preprocess your data
5 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2
6
7 # Create and train a multi-class classification model (e.g., support ve
8 model = SVC()
9 model.fit(X_train, y_train)
10
11 # Make predictions
12 predictions = model.predict(X_test)
13 predictions
```

```
Out[62]: array([0, 0, 2, 2, 0, 2, 0, 0, 1, 1, 1, 0, 1, 0, 1, 1, 2, 1, 2, 0, 0, 0,
2, 0, 0, 1, 1, 2, 2, 1])
```

```
In [48]: 1 import numpy as np
2 from sklearn import datasets
3 from sklearn.model_selection import train_test_split
4 from sklearn.svm import SVC
5 from sklearn.metrics import accuracy_score, classification_report
6
7 # Load the Iris dataset
8 iris = datasets.load_iris()
9 X = iris.data
10 y = iris.target
11
12 # Split the data into training and testing sets
13 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3
14
15 # Evaluate the classifier's performance
16 accuracy = accuracy_score(y_test, y_pred)
17 report = classification_report(y_test, y_pred, target_names=iris.target
18 # Print the results
19 print("Accuracy:", accuracy)
20 print("Classification Report:\n", report)
21
```

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	19
versicolor	1.00	1.00	1.00	13
virginica	1.00	1.00	1.00	13
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45

multi-label classification

```
In [60]: 1 from sklearn import datasets
2 from sklearn.model_selection import train_test_split
3 from sklearn.svm import SVC
4 from sklearn.metrics import accuracy_score
5
6 # Load a sample dataset (Iris dataset as an example)
7 iris = datasets.load_iris()
8 X = iris.data
9 y = iris.target
10
11 # Split the data into training and testing sets
12 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
13
14 # Create a Support Vector Machine (SVM) classifier
15 clf = SVC(kernel='linear', C=1)
16
17 # Fit the classifier to the training data
18 clf.fit(X_train, y_train)
19
20 # Make predictions on the test data
21 y_pred = clf.predict
22 y_pred
```

Out[60]: <bound method BaseSVC.predict of SVC(C=1, kernel='linear')>

Confusion Matix

```
In [34]: 1 # Import the necessary Libraries
2 from sklearn.metrics import confusion_matrix
3 import numpy as np
4
5 # Actual and predicted Labels (replace these with your actual data)
6 actual = np.array([1, 0, 1, 1, 0, 1, 0, 1, 0, 0])
7 predicted = np.array([1, 0, 1, 0, 1, 1, 0, 1, 0, 1])
8
9 # Create the confusion matrix
10 cm = confusion_matrix(actual, predicted)
11
12 # Print the confusion matrix
13 print("Confusion Matrix:")
14 print(cm)
15
```

Confusion Matrix:

```
[[3 2]
 [1 4]]
```

```
In [33]: 1 import numpy as np
2
3 # Given confusion matrix
4 confusion_matrix = np.array([
5     [13, 45],
6     [0, 15]
7 ])
8
9 # (i) Accuracy
10 total_samples = np.sum(confusion_matrix)
11 correct_predictions = np.trace(confusion_matrix)
12 accuracy = correct_predictions / total_samples
13 print(f"(i) Accuracy: {accuracy:.2f}")
14
15 # (ii) Precision for Class 1
16 precision_class_1 = confusion_matrix[1, 1] / (confusion_matrix[0, 1] +
17 print(f"(ii) Precision for Class 1: {precision_class_1:.2f}")
18
19 # (iii) Recall for Class 1
20 recall_class_1 = confusion_matrix[1, 1] / (confusion_matrix[1, 0] + con
21 print(f"(iii) Recall for Class 1: {recall_class_1:.2f}")
22
23 # (iv) Specificity for Class 0
24 specificity_class_0 = confusion_matrix[0, 0] / (confusion_matrix[0, 0]
25 print(f"(iv) Specificity for Class 0: {specificity_class_0:.2f}")
26
27 # (v) F1-Score for Class 1
28 f1_score_class_1 = 2 * (precision_class_1 * recall_class_1) / (precision
29 print(f"(v) F1-Score for Class 1: {f1_score_class_1:.2f}")
30
```

(i) Accuracy: 0.38
(ii) Precision for Class 1: 0.25
(iii) Recall for Class 1: 1.00
(iv) Specificity for Class 0: 0.22
(v) F1-Score for Class 1: 0.40

decision tree

```
In [58]: 1 # Import the necessary libraries
2 from sklearn import tree
3
4 # Create a dataset for training the decision tree
5 # In this example, we have two features (X) and a target variable (y)
6 X = [[0, 0], [1, 1]]
7 y = [0, 1]
8
9 # Create a decision tree classifier
10 clf = tree.DecisionTreeClassifier()
11
12 # Train the classifier on the dataset
13 clf = clf.fit(X, y)
14
15 # Make predictions using the trained decision tree
16 predictions = clf.predict([[2, 2], [0.5, 0.5]])
17
18 # Print the predictions
19 print(predictions)
20
```

[1 0]

decision tree classifier

```
In [59]: 1 from sklearn import datasets
2 from sklearn.model_selection import train_test_split
3 from sklearn.tree import DecisionTreeClassifier
4 from sklearn.metrics import accuracy_score
5
6 # Load a sample dataset (Iris dataset as an example)
7 iris = datasets.load_iris()
8 X = iris.data
9 y = iris.target
10
11 # Split the data into training and testing sets
12 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
13
14 # Create a Decision Tree Classifier
15 clf = DecisionTreeClassifier()
16
17 # Fit the classifier to the training data
18 clf.fit(X_train, y_train)
19
20 # Make predictions on the test data
21 y_pred = clf.predict(X_test)
22
23 # Calculate the accuracy of the classifier
24 accuracy = accuracy_score(y_test, y_pred)
25 print(f"Accuracy: {accuracy * 100:.2f}%")
26
```

Accuracy: 100.00%

In []:

1